CHRISTIAN-ALBRECHTS-UNIVERSITY

MASTER THESIS

# Alternative Software Transaction Implementation in Haskell

*Author:*
Lasse Folger

*Supervisor:*
Dr. Frank Huch

*A thesis submitted in fulfillment of the requirements
for the degree of Master of Science*

*in the*

Programming Languages and Compiler Construction
Department of Computer Science

April 13, 2017

# Declaration of Authorship

Hiermit erkläre ich an Eides statt, dass ich die vorliegende Arbeit selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel verwendet habe.

Kiel, April 13, 2017

_____

Christian-Albrechts-University

# *Abstract*

Faculty of Engineering
Department of Computer Science

Master of Science

**Alternative Software Transaction Implementation in Haskell**

by Lasse Folger

Multi core processor architecture are omnipresent today. Even mobile devices contain multi core processors. To ultilize multicore processors to their full extend it is needed to write multithreaded code. This on the other hand requires synchronization between the threads. While lock based synchronization concepts yield the problem of deadlocks, are lock-free primitives often hard to use. Thus software transactions were developed, a way to avoid both of these problems for a drawback in performance. Haskell also offers an implementation for software transactions, called STM. It is implemented as a Haskell library that makes use of C primitives and provides a clear interface. In this thesis I analyze the problems that lead to unnecessary rollbacks. Furthermore do I provide a sulotion to avoid these problems and present an alternative pure Haskell implementation that offers the same interface as the original library. The performance of the different implementations is compared and the results are presented here as well. These results show the success of the new implementation, even though it does not make use of C primitives to improve its performance.

# Contents

viii

# List of Figures

# List of Abbreviations

| | |
|---|---|
| **OS** | **O**perating **S**ystem |
| **CPU** | **C**entral **P**rocessing **U**nit |
| **WHNF** | **W**eak **H**ead **N**ormal **F**orm |
| **STM** | **S**oftware **T**ransactional **M**emory |
| **ACID** | **A**tomicity **C**onsistency **I**solation **D**urability |
| **TVar** | **T**ransactional **Var**iable |
| **MVar** | **M**utable **Var**iable |

# Chapter 1

# Motivation

Modern computer architecture includes multicore processors. To utilize these multicore system to their full extend, concurrent and parallel programming is needed. By this new challenges arise. There are technical challenges. For example a new scheduler is needed and hardware accesses (Printer, Display, etc.) need to be coordinated. These are challenges the operating system usually handles. There are other challenges the operating system cannot handle, because they are specific for every program.

One challenge is the splitting of the problem in smaller problems which can be processed by different threads in parallel. This challenge is very specific to each problem and there is no general solution for it. The most discussed challenge is the synchronization. If a program works with multiple threads, these threads usually communicate. Communications means to exchange data. Even a simple statement like an assignment can cause problems when used in the parallel threads. The problem is that these operations are non atomic operations. Thus ($x$ = $x$ + $1$) consist of three parts. first reading the old value, second adding $1$, and thrid write the new value. This means two threads in parallel can both read the old value, then both add $1$ to the old value, and then write back the new value. The new value is the initial value incremented by $1$, even though two threads executed an increment operation on this value. This non inteded behaviour is called *lost update*. The efforts to avoid such non intended behaviour are called synchronization.

Although multicore processors are new, the research in the field of synchronization has a long history, starting with (Dijkstra, 1965), who introduced the most basic synchronization tool, the semaphore. The semaphore is a abstract datatype which holds an Integer as state and provides two *atomic* operations, P and V. If the value of the semaphore is greater than $0$, P decrements the semaphore. If the value of the semaphore is $0$, the thread that evoked P is suspended. When a thread evokes V the value of the semaphore is increased and in the case another thread is currently suspended, because it called P on the semaphore, that thread resumes. When the thread resumes, it tries P again.

This seem to be a simple construct, but its capabilities are enormous. It is highly complex to use a semaphore correctly. The main problem of semaphores is the so called deadlock[1]. This means there is a schedule, where no progress of the systen is possible, because all threads are waiting for a semaphore. The term deadlock is not exclusive for semaphores. It is used for all blocking mechanisms. Avoiding deadlocks is very hard even when using one or few semaphores. It is nearly impossible to avoid deadlocks when you try to compose semaphore based functions.

To avoid the problems of semaphores while maintaining the expressiveness of semaphores in Haskell, software transactions were introduced (Harris, Marlow, et

---

[1]In the course of this thesis I will refer to deadlocks as a static propertie rather than a state of a system.

al., 2005). Sofware transactions are inspired by the long known database transactions (Gray and Reuter, 1992). Software transactions in Haskell provide an interface to program with single element buffers. If you are using this interface the underlying implementation ensures the *ACI(D)* properties.

There is a stable implementation for software transactions in Haskell, namely Software Transactional Memory (called STM in the following). The STM library provides an interface which allows the user to process arbitrary operation on one element buffers. The operations can be grouped to *transactions*. When a transaction is executed the library ensure the ACI(D) properties. This is done by optimistically executing the transaction. If a conflict is detected, the changes of the transaction are discarded and the transaction is restarted (also called rollback). This works, but is not optimal with regards to efficiency and performance. There are two problems. First the conflict detection. Sometimes the implementation detects a conflict and evokes a rollback, even though it is not necessary. The second problem is the rollback mechanism. Regardless the conflict, always the whole transaction is reexecuted. This includes operations on data that has not changed, thus an unnecessary recomputation. These problems are discussed in detail in Chapter 2. The aim of this thesis is to provide an alternative implementation that avoids these problems while preserving the ACI(D) properties. Chapter 3 provides concept to solve one of these problems, while Chapter 4 discusses the implementation of this concept. The performance of the alternative implementation is examined in Chapter 5. The last Chapter provides an overview of related work and gives several directions for improving the alternative implementation. In the final part of the last Chapter the thesis is summarized and concluded.

# Chapter 2

# Introduction

## 2.1 Software Transactional Memory

Software Transactional Memory (STM in the following) is a programming language independent synchronization concept. Today STM is available in all common programming languages. Since the subject of this thesis is Haskell, we will not investigate STM in general. To understand the benefits of STM, take a look at the following example:

```haskell
type Account = MVar Int

transfer :: Account -> Account -> Int -> IO ()
transfer src dst am = do
  balSrc <- takeMVar src
  balDst <- takeMVar dst
  putMVar src (balSrc - am)
  putMVar dst (balDst + am)
```

This is a simple implementation of a bank account and an associated transfer function. This implementation uses an `MVar` for synchronization. An `MVar` is a buffer with a capacity of one. This buffer can either be empty or filled. If the MVar is empty, every `takeMVar` operation on this MVar blocks until the MVar is filled. If the MVar is filled, `takeMVar` empties the `MVar` and return the value. `putMVar` is the opposite operation. It fills the MVar with a value, if it is empty and suspends if the MVar is already filled.

This means `transfer` first empties both `Accounts`, then modifies the balances and at last writes back the new balances. At first glance this function seems to work fine, but the following example contains a deadlock:

Thread 1:

```
main = do
  transfer acc1 acc2 50
```

Thread 2:

```
main = do
  transfer acc2 acc1 50
```

The problem is the mutual access of the MVars. If both threads take their `src` at the same time, they will both wait for `dst` [1]. To avoid this deadlock we can rewrite the code:

```haskell
transfer src dst am = do
  srcBal <- takeMVar src
  putMVar src (srcBal - am)
  dstBal <- takeMVar dst
  putMVar dst (dstBal + am)
```

---

[1] In fact, `transfer acc1 acc1 50` is enough to evoke a deadlock

This indeed solves the problem regarding the deadlock. In return we lose consistency. For a brief moment an we see an inconsistent state. The amount is allready withdrawn from one account, but not yet deposited on the other account. This inconsistent state is observable by other threads. It is not possible in the first implementation.

We can use STM to avoid both of these problems. STM provides a single element buffer named `TVar`. In contrast to an MVar, a TVar always holds a value and is never empty. TVars are read and written with the functions `readTVar` and `writeTVar`, respectively. In contrast to `putMVar` and `takeMVar`, the TVar operations are not `IO` actions but `STM` action[2]. `STM` is an instance of Monad, hence multiple `STM` actions can be combined using the comfortable do-notation. The following code represents the example from above; implemented with TVars instead of MVars:

```
type Account = TVar Int

transfer :: Account -> Account -> Int -> STM ()
transfer src dst am = do
  srcBal <- readTVar src
  dstBal <- readTVar dst
  writeTVar src (srcBal - am)
  writeTVar dst (dstBal + am)
```

Note that `transfer` is no longer an `IO` action, but an `STM` action. Apart from this the code looks similar to the MVar version.

The function `atomically ::  STM a -> IO a` is used in order to execute an STM action. The following example contains no deadlock:

Thread 1:                                        Thread 2:

```
main = do
  atomically $
    transfer acc1 acc2 50
```

```
main = do
  atomically $
    transfer acc2 acc1 50
```

This is because STM ensures the *ACID* properties. The ACID properties were Introduced in (Gray and Reuter, 1992) for database transactions. These properties were adapted for software transactions later on. In the case of software transactions the ACID properties mean the following:

- *Atomicity*: all operations of the transaction are executed or none.

- *Consistency*: all modifications of a transaction are committed at the same time. No transition state is observable.

- *Isolation*: no concurrency is observable by any transaction. Transactions do not influence each other indirectly.

- *Durability*: ensures the perseverance of the changes.

In the case of software transactions the *Durability* is not demanded, which is why we will refer to the ACI properties for the rest of the thesis.

These properties explain the name `atomically`. The enclosed code appears to be executed instantaneously without any interactions with other threads. In the following any composition of `STM` actions is called *transaction*. Before we turn over to the implementation of STM, we take a deeper look at the interface of the STM.

---

[2]If you are wondering when I use `SMT` and when STM. I use `STM` when I refer to the Haskell type constructor and STM when I refer to STM as library

`newTVar ::   a → STM (TVar a)` creates a newTVar. Since a TVar always holds a value, an initial value has to be passed to create a TVar. There is no function like `newEmptyTVar`.

Besides functions to create and access TVars, there are functions to alter the control flow. `retry ::   STM a` is a generic STM action that indicates a failure, thus whenever a transaction engages a `retry` it restarts. The transaction is **not** restarted immediately. The transaction restarts when at least one of the TVars it has read is modified. If the transaction restart immediately (and no TVar has changed), the transaction executes the same code, including `retry`.

With `orElse ::   STM a → STM a → STM a` you are able to express alternatives. `orElse` executes the first transaction and ignores the second transaction if the first transaction is successful. If the first transaction fails (retries), the second transaction is executed instead.

Note that it is not possible to execute `IO` action within a transaction, which means that no side effects can occur. Furthermore this means restarting a transaction will never lead to the re-execution of irreversible operations. The reason is that the computations of transactions are performed within the `STM` monad. In other words, the type system of Haskell forces us to write correct transactions.

Abstraction and composability are key features for single threaded programming. These features allow us to combine small pieces of code to complex functions. They are not available for lock based programming. Composing correct lock based functions often leads to deadlocks or inconsistencies. Consider the following example:

```
withdraw ::  Account -> Int -> IO()
withdraw acc am = do
  bal <- takeMVar acc
  putMVar acc (bal - am)

deposit ::  Account -> Int -> IO()
deposit acc am = withdraw acc (-am)
```

These are functions to withdraw and deposit money from and on an account. The natural way to implement `transfer` is:

```
transfer ::  Account -> Account -> Int -> IO()
transfer src dst am = do
  withdraw src am
  deposit dst am
```

We reuse the functions that are already defined instead of coding everything from the scratch. In our example this is equivalent to the solution suggested above to eliminate the deadlock. This implementation is free of deadlocks, but it lacks consistency. Thus building complex concurrent operations can not take advantage of abstraction and composability. We always need to code everything from the scratch. This is error prone in comparison to the step wise combination of smaller operations into more complex operations.

STM allows us to use this important programming paradigm for concurrent programming. Thus the following example provides deadlock freedom as well as consistency.

```
withdraw ::  Account -> Int -> STM()
withdraw acc am = do
  bal <- readTVar acc
```

```
  writeTVar acc (bal − am)

deposit :: Account −> Int −> STM()
deposit acc am = withdraw acc (−am)

transfer :: Account −> Account −> Int −> STM()
transfer src dst am = do
  withdraw src am
  deposit dst am
```

We can combine arbitrary transactions to more complex transactions while preserving the ACI properties. This greatly benefits the readability of the code. In addition it increases the efficiency of the development process because we are able to reuse code that was already found to be correct. This was also one of the main motivations for the paper of (Harris, Marlow, et al., 2005) which forms the foundation of STM in Haskell.

The reason this works is because always the whole transaction is considered as one block for which the ACI properties must hold. Thus the user marks the critical section by defining it as one transaction and the library ensures the correctness and deadlock freedom. The user only needs to think about which actions have to be processed together. This is comparable to a lock based version with a single lock. Everytime the user wants to process a critical section he takes the lock before this section and releases the lock afterwards. Then the critical section is processed isolated and problems such as race conditions and lost updates do not occur. The performance on the other hand is devestating and does not scale well, because all critical sections are sequentialized. For most modern this is systems not acceptable, thus this solution is not feasible.

## 2.2   Implementation

In this section we explore the current implementation of STM in Haskell, more specific in GHC.[3] Even though the current implementation uses low level C primitives, we retain an abstract view on the implementation, since the technical details are not important for the course of this thesis. The implementation is outlined to understand how the ACI properties are guaranteed.

The execution of a transaction (a call of `atomically`) is split in two phases. First the computation phase and second the commit phase.

### 2.2.1   Computation Phase

Each transaction holds a log for the TVars that it has accessed. The log contains four elements per entry. These are:

- `tvar`

- `expectedValue`

- `newValue`

- `versionNumber`

---

[3]For a detailed description of the implementation refer to `https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/STM`.

The `versionNumber` is only used to prevent a very subtle bug and thus not considered in this thesis. The log is extended and modified by the transactional operations `writeTVar` and `readTVar`. `newTVar` on the other hand creates the new TVar directly.

Whenever `readTVar` is called the associated TVar is looked up in the log. If it is present, the `newValue` is returned. If it is not present, a new entry in the log is created. While `tvar` is the passed TVar, `newValue` and `expectedValue` are the current value of the TVar. This is one of the two times in the computation phase when the transaction accesses the actual mutable data structures. After the entry is created and added to the log, the current value is returned.

A call of `writeTVar` also looks up the associated TVar in the log. If it is present, the field `newValue` is set to the value passed to `writeTVar`. If it is not present, a new entry is created. The `tvar` is the passed TVar, the `newValue` is the passed value and `expectedValue` is the current value of that TVar. This is the other time the actual mutable data structures are accessed in the computation phase. Entering the current value of the TVar at this point is not needed to preserve the ACI properties. This is done to simplify the implementation at the risk of an additional rollback. On the other hand it is unusual to write a TVar that was not read before, because this means overwriting the value of a TVar and thus discarding the initial content of that TVar.

This log fulfills two purposes. One purpose of the log is the use in the commit phase which is described in Section 2.2.2. The other is the interaction between `readTVar` and `writeTVar`. The `readTVar` operations are able to use the results of preceding `writeTVar` operations. Without the log `writeTVar` would need to access the actual TVar. This on the other hand implies that other transactions are able to see inconsistent intermediate states of the system; a violation of the ACI properties. It may seem unnecessary to read a TVar that the transaction itself wrote before. The transaction should know what it writes and thus does not need to access such TVars. However, there are two reasons to allow it. The current implementation allows the user to combine all transactional actions in an arbitrary manner and the library ensures (at compile time) that it works correctly. To restrict the user to only read TVars that he has not yet written, no longer allows the library to give this kind of guarantee at compile time; this contradicts the design concept of Haskell. The second reason is one of the core motivations of STM in Haskell: composability. With the restriction it is not possible to combine arbitrary correct STM functions to new more complex STM functions.

To understand how the log is constructed, take a look at the following example:

```
transaction = do
  a <- readTVar t1
  b <- readTVar t2
  writeTVar t1 b
  writeTVar t2 a
```

This code would lead to the following log:

```
log  = {(t1,a,b),(t2,b,a)}
```

The log contains two entries because the transaction accessed two TVars. The first part of the entry denotes the TVar, the second part the expected value and the last part is the new value. The first entry contains the information that `t1` held the value `a` when it was first read and the value `b` is the new value of it. `t2` held the value `b` and the new value is `a`. Before we will examine the commit phase, we will look at the other operations of the STM interface.

`newTVar` creates a new TVar and initializes this TVar. Afterwards this TVar can be used like already existing TVars. Even if the transaction is rolled back, the new created TVars are not deleted explicitly. This work is done by the garbage collector since the TVars are not further referenced.

`retry` aborts the computation and returns a results that indicates a failure. This result may be intercepted by `orElse` or is passed to `atomically` directly.

If `atomically` receives an result that indicates an failure, it aborts the transactions. Aborting a transaction means to discard the log. Since no observable operations are performed in the computation phase, nothing has to be undone. As soon as at least one of the TVars in the log has changed, the transaction is restarted. If the transaction is restarted immediately and no TVar has changed the transaction would reach the same `retry` again. These changes can be checked by comparing the `expectedValue` in the log with the current value in the TVar. To avoid busy waiting the thread does not repeatedly check if the value has changed. The TVar has a queue for waiting threads. Each time a transaction successfully commits and writes a TVar it also checks if there is someone waiting in this queue. The committing thread then notifies all waiting threads.

`orElse` on the other hand reacts differently on the the result that indicates a failure. The implementation works with nested transactions, but to explain this in detail would go beyond the scope of this thesis. Nested transactions are not able to publish their writes on their own. When a nested transaction successfully commits, its log is integrated in the log of the surrounding transaction. Integrated means the logs are merged and in the case that there is a entry in both logs for one TVar, the entry of the outer transaction is discarded. If the nested transaction fails, because `retry` occurred and it is the first transaction of `orElse`, the log of the inner transaction is integrated in the log of the surrounding transaction, but the `newValue` fields of the inner log are ignored. Afterwards the second transaction is executed with this log. If a nested transaction fails to validate (explained in the next section) the outermost transaction is rolled back.

In conclusion the interface functions of STM are processed in the computation phase as follows:

- `writeTVar`: Look up TVar in log. If present update `newValue`. If not present read actual TVar and create new entry.

- `readTVar`: Look up TVar in log. If present return `newValue`. If not present read actual TVar and create new entry.

- `newTVar`: Create and initialize a new TVar.

- `retry`: Return a result that indicates a failure.

- `orElse`: Create a nested Transaction and reacts on the return value of that transaction.

### 2.2.2   Commit Phase

After the log is calculated and no further STM actions need to be processed, the commit phase starts. At first the transaction checks if the values in its log are still correct by *validating* its log. Validation denotes the process to check if the `expectedValues` are equal to the current values in the TVars. In other words, for each entry in the log, the transaction compares the `expectedValue` in the log with the current value of

the actual TVar. If at least one of these values does not match, the transaction is considered *invalid*. If the validation returns invalid the transaction is instantaneously rolled back. If all values match the the transaction is considered *valid*. If the validation returns valid, each entry in the log is processed again.

If `expectedValue` differs from `newValue` the associated TVar is locked. When the transaction has acquired all locks that it needs, it validates again. This seems wasteful in terms of resources, but locking the TVars is considered an expensive operation and thus the implementation tries to avoid this whenever possible. This double validation eliminates the chance that the transaction acquires all locks and then finds out it is invalid; consequently a unnecessary locking of TVars. If the validation fails at this point, the transaction is rolled back after the locks has been released. If the transaction has acquired all locks and is valid the transaction is is ready to publish its changes. The transaction iterates through the log and updates the actual TVars if `expectedValue` and `newValue` differ. Simultaneously the transaction releases the locks for the TVars.

If the validation returns invalid it means at least one `expectedValue` is no longer correct. To roll back is essential to retain the ACI properties. The failed validation indicates that transaction has read a value that is now outdated and possibly worked with this value. Take a look at the following example:

```
transaction = do
  a <- readTVar t1
  writeTVar t1 (a+1)
```

If this transaction is processed by two transactions in parallel. Both would read the initial value of `t1`, say 1. So both would note in their log `(t1,1,1)`. After the `writeTVar` the log of both transactions contains `(t1,1,2)`. After that both transactions try to commit. Assume one transaction commits before the other transaction tries. [4] Then the transaction would find its log to be valid and lock `t1`. After that its log is still valid and so it modifies `t1` and releases the log. Then the second transaction tries to commit. Since the actual value hast changed to 2 it does no longer match the `expectedValue` and the transaction is rolled back. If the transaction would not be rolled back at this point and commit instead. The transaction would write 2 to the TVar (that already contains 2). In the end this means the value of `t1` is 2 after both transactions have finished. This is certainly not the intended behaviour after incrementing the TVar twice that holds 1. This is the well known *lost update* problem. By rolling back the second transaction it reads `t1` once more. The log contains `(t1,2,2)` after the `readTVar` operation and `(t1,2,3)` after the `writeTVar` operation. Then the transaction validates, locks, validates and finally publishes it modifications. In the end the value of `t1` is 3; just as intended.

### 2.2.3 Notes on the Implementation

(Larus and Rajwar, 2007, Chapter 2) describe in their book different design options of the implementation of a (Software) Transactional Memory. While most of these options effect only the performance of a system, some also effect the semantics of the system. We will discuss in this section the design options that are important for this thesis [5].

---

[4]For simplicity we assume that no other transaction is running besides the two we are looking at.

[5]The names used in the following part are taken from (Larus and Rajwar, 2007, Chapter 2)

**Deferred and Direct Updates**

The way a STM system modifies the underlying data structures can either be *deferred*
or *direct*. Direct updating systems are writing the actual objects when a write oper-
ation is called. In the case of Haskell this would mean, every time `writeTVar` is
called. Deferred updating systems on the other hand buffer the write operations to
commit them later on. Haskell STM is a deferred updating system since the values
are buffered in the log before they are committed. This design options does not ef-
fect the semantics of the system. While a direct system loses performance, when a
transaction is rolled back because the initial values need to be restored, a deferred
systems contains an overhead due to the need to log values and looking them up.
Neither mechanism is better than the other in general; it depends on the application
that STM is used in. (Harris, Plesko, et al., 2006) compare a deferred and a direct
system. They show that the performance of a direct update system is significantly
higher than that of a deferred system, when reads outnumber writes by far.

**Early and Late Conflict Detection**

A STM system needs to detect conflicts in order to ensure the ACI properties. This
can either be done as soon as the conflict occurs or later before the transaction com-
mits. If the system uses a late conflict detection, transactions may work on an in-
consistent state. This may lead to loops or exceptions. So this design decision is
relevant for the semantics. In Haskell, STM uses a late conflict detection. A possible
conflict is detected by validating the log before comitting the transaction. This im-
plies the transaction may work on an inconsistent state until it attempts to commit.
This means the transaction may run into an infinite loop because it saw an inconsis-
tent state. To avoid this problem, additional validations are performed each time the
executing thread yields. Exceptions raised in the transaction are handled as follows.
If the log is valid, the transcation raises the exception to the caller. If it is invalid the
transaction is restarted immediately. In conclusion, the user of STM can not observe
that the transaction worked on an inconsistent state.

**Synchronization**

Another important property of a STM system is the way it synchronizes transac-
tions. In order to validate correctly the systems needs to make sure the validation
result does not depend on race conditions and is correct until the commit is com-
pleted. This means either concurrent transactions are delayed or their commits do
not change each others validity. In Haskell the first approach is taken. When a
transaction commits, the TVars in the log that are updated are locked, thus other
transactions that may conflict are not able to commit at the same time. In order to
avoid a deadlock, all locks are released and the transaction is rolled back when it
tries to aquire the lock for a locked TVar. In the worst case this leads to the roll back
of both transactions, however the chances are narrow. Rolling back the transaction
seems to be harsh instead of waiting until the other transaction finishes and then
trying to commit, but if two transactions try to lock the same TVar, both transactions
try to write this TVar. This means at least one of the transactions is rolled back, since
the TVar is logged with the old value. Thus the first transaction to commit would
modify a value. The other transaction then sees this update and rolls back.

**Strong and Weak Isolation**

When executing multiple transactions, STM must be able to isolate these transactions. The ACI properties demand this. What happens if non transactional code work on the same shared data structures is up to the STM implementation. If it isolates transactional and non transactional code, it is called strong isolation; otherwise, weak isolation. The STM library has limited possibilities to effect non transactional code. It is not possible to roll back such code. Either STM prevents external accesses completely or it suspends external accesses until no transaction is using the shared data. STM in Haskell uses the first approach. The only way to access TVars is to use the functions `readTVar` and `writeTVar`. The type system ensures that every access to TVars is in an `STM` action. The only way to execute `STM` actions on the other hand is to use `atomically`, by which this action becomes a transaction. STM in Haskell provides strong isolation although it is possible to read TVars with `readTVarIO`. As the name suggests this function returns an IO action which reads a single TVar. However, it is not possible to modify transactional variables with IO actions.

## 2.3 Problems

In this section we turn over to the problems in the current implementation. These problems can be examined independently. The first problem is about *when* a transaction is rolled back and the second problem is about *how* a transaction is rolled back.

### 2.3.1 Unnecessary Rollback

Remember the STM implementation of `transfer` and its example use given in 2.1:

```
transaction1 = do
  atomically $
    transfer acc1 acc2 50
```

```
transcation2 = do
  atomically $
    transfer acc2 acc1 50
```

The implementation is correct, but not very efficient in this case. Take a look at the inlined functions to understand the problem:

```
transaction1 = do
  a1 <- readTVar acc1
  a2 <- readTVar acc2
  writeTVar acc1 (a1 - 50)
  writeTVar acc2 (a2 + 50)
```

```
transaction2 = do
  a1 <- readTVar acc2
  a2 <- readTVar acc1
  writeTVar acc2 (a1 - 50)
  writeTVar acc1 (a2 + 50)
```

Due to the scheduler the threads can run in a sequential order. This case may occur, but is not desirable. It means there is no performce improvement by executing this on multiple cores. Thus the efforts to use multiple threads are futile in the first place. This is not a problem specific to STM, but to all synchronization mechanisms. If the resulting multi threaded program is not scheduled in a way that it is executed parallel, these mechanisms are a performance deterioration rather than a performance improvement. Since we cannot access the scheduler, we ignore this case. Although we sometimes use concurrent programming to enable reactivity of a program, performance is always an important factor. A bad performance results in a bad reactivity.

FIGURE 2.1: Time when the update of `acc1` or `acc2` causes a roll-back.

The second case is that these transactions are run in parallel. This should be the better case, because the implementation has a chance to improve the performance. Unfortunately this is not the case. To understand why, we need to take a close look at the execution. Let us assume both threads execute their computation phase at the same time. This means both read the initial values of `acc1` and `acc2` and add these information to their logs. Furthermore add both transactions entries for `writeTVar acc1` and `writeTVar acc2` to their logs. Then both transactions try to commit, thus try to lock the TVars. It is even possible that both transactions are rolled back at this point. Lets assume `transaction1` acquires the locks for `acc1` and `acc2`. Since no TVars were modified after `transaction1` read them, it validates and commits. If `transaction2` tries to access the TVars before `transaction1` has finished committing, it is rolled back. Thus it is possible for `transaction2` to read the old value once again. If `transaction2` is descheduled for the time `transaction1` commits, it is rolled back afterwards, because the values of `acc1` and `acc2` have changed by `transaction1`. In conclusion no performance improvement was achieved. The most efficient execution is if both transactions are executed in a sequential order. As mentioned before, this not desirable for multithreaded programs.

This leads to two questions:

- When is it needed to roll back a transaction?

- How can we avoid or at least decrease rollbacks?

A transaction needs to be rolled back if it is operating on data that is not a snapshot of the current memory. In other words if a value has changed after the transaction read this value. When a transaction reads a TVar, this TVar becomes *critical* for the transaction (see Section 2.4 for definition). Critical means a modifications of that TVar causes the transaction to roll back. Figure 2.1 visualizes when the TVars `acc1` and `acc2` are critical for `transaction1`. When `readTVar acc1` is executed the values becomes critical and stays critical until the transaction commits. If any other transaction commits a modification to `acc1` or `acc2`, while `acc1` and `acc2` are critical for `transaction1`, `transaction1` is rolled back to preserve the ACI properties.

This insight brings an intuitive way to deal with this problem. If we minimize the time the TVars are critical for a transaction, we reduce the chance that this transaction is rolled back. If we rearrange the operations of `transfer`, we are able reduce the time `dst` is critical. Note that we can rearrange the operations to a certain degree without changing the semantics of the resulting code due to the ACI properties.

```
transfer src dst am = do
```

FIGURE 2.2: Effect of rearranging code with regards to the time `acc1` and `acc2` are critical for `transaction1`.

```
srcBal <- readTVar src
writeTVar src (srcBal - am)
dstBal <- readTVar dst
writeTVar dst (dstBal - am)
```

With this implementation of transfer the time in that both TVars are critical is reduced. Figure 2.2 shows the effects of this for `transfer acc1 acc2 50`. The second TVar, namely `acc2`, is shorter critical than in the initial implementation. The time `acc1` is critical has not changed at all. Nevertheless, it shows that delaying the execution of `readTVar` can reduce the time TVars are critical and by this the chance the transaction is rolled back. Our aim is to delay the execution of `readTVar` as far as possible to reduce the time a TVar is critical for a transaction. We have already seen one option to achieve this; rearrange the operations of the transaction. To apply this in general, it would require a kind of preprocessing in the compiling process, for example a source to source code transformation. The aim of this thesis is to provide an pure Haskell library. I do not intend to implement an extension to the compiler nor do I want to provide a source to source code transformer. The only other option is to alter implementation of `readTVar` and `writeTVar` without changing the *external* semantics of STM. External semantics are the semantics the user can observe and which effect the user written code.

The critical time would be minimal if the TVars were read directly before or at the start of the commit phase. Then the chances that another transaction commits a change to a TVar that is critical are low or non existing. So the idea is to let the user define transactions like before, but changing the semantics of `readTVar` that it is evaluated in the commit phase. As a result, the user is able to define transactions like in the original implementation, but the delay of the evaluation shortens the time TVars are critical and thus the chances a transaction is rolled back.

Lets refer to our example:

Thread 1:

```
transaction1 = do
  a1 <- readTVar acc1
  a2 <- readTVar acc2
  writeTVar acc1 (a1 - 50)
  writeTVar acc2 (a2 + 50)
```

Thread 2:

```
transaction2 = do
  a1 <- readTVar acc2
  a2 <- readTVar acc1
  writeTVar acc2 (a1 - 50)
  writeTVar acc1 (a2 + 50)
```

If we change the semantics of `readTVar` by delaying the evaluation, the following happens. Both transactions will execute the computation phase simultaneously. This means `transaction1` adds `(acc1,a1,(a1 - 50))` and `(acc2,a2,(a2 +`

50)) to its log (this is analog for `transaction2`). At the first glance this seems to be incorrect since the values of `a1` and `a2` are not yet present. For Haskell this is quite common. Haskell is a non-strict language which means passing unevaluated expressions is normal.

After the computation phase the commit phase follows. The first step is to lock the read TVars in order to perform the validation. Since both transactions used the same TVars, they will commit successively instead of parallel.

Assume `transaction1` gets the locks first and tries to validate[6]. Since the log contains an action whose result is the current value, the validation is unnecessary; it is always valid. To validate the log, `a` and `b` are evaluated. At last the new values are written to the TVars.

After `transaction1` finished and released the locks, `transaction2` aquires these locks and validates. The log of `transaction2` is also valid and also commits its changes.

Both transactions run parallel as far as possible and did not roll back. Chapter 3 presents the limitations of this idea and the challenges that arise when implementing it.

### 2.3.2 Unnecessary Recomputations

While the first problem dealt with then question *when* transactions need to be rolled back, the second problem investigates the question *how* transactions are rolled back. Lets take a look at our well known example:

```
transfer src dst am = do
   srcBal <- readTVar src
   writeTVar src (srcBal - am)
   dstBal <- readTVar dst
   writeTVar dst (dstBal + am)
```

This transaction contains two independent statements. The first two lines of the transacton form the first statement. This is independent of the last two lines. Independent means their side effects or results do not influence each other. While the first line influces the second line, it does not influence the last two lines and vice versa.

If the transaction is executed, it computes its log first. Then it validates and locks the TVars [7]. The validation fails if either of the TVar has changed after it was read by the transaction. If the validation fails the transcation is rolled back. Which means the log is discarded, regardless which TVar was the reason for the failed validation.

Suppose a `transaction1` executes `transfer acc1 acc2 5` and is deschedules before committing. Then `transaction2` modifies `acc1` (and nothing else) and commits. This would cause `transaction1` to roll back and execute both parts of `transfer` again. This includes the read and write of `t2`, although `t2` was not modified. Hence the exact same code with same inputs and the same (relevant) environment is executed twice. If we just execute the parts of a transaction that are invalid instead of all, we can save a considerable amount of time when a transaction is rolled back. This means for the example, it is enough to remove the entry for `acc1` from the log and execute the first two actions of transfer instead of all actions.

---

[6]You could argue that evaluating `readTVar` operation is necessary before validating, but this would not change the validity of the transaction, since the TVars are locked and can not be modified by other transactions at that point.

[7]We want study the two problems independently and thus assume the original implementation here.

## 2.4 Terminology and Conventions

To avoid misunderstandings, we use this section to define the meaning of specific terms and explain code conventions.

**Conflict**

A *conflict* in STM is if two transactions access the same TVar in the following way. Transaction `t1` reads the TVar. After `t1` has read the TVar and before it has commited, transaction `t2` writes this TVar (calls `writeTVar`). This means that `t1` works with an outdated value after `t2` commits. In the current implementation conflicts are not detected until both, `t1` and `t2`, try to commit. If only `t1` commits, it does not know that `t2` has written the TVar. If only `t2` commits, it does not know that `t1` has read this TVar. This allows both transactions to commit without rolling back either of them, if `t1` commit before `t2` (and `t1` does not write any TVar that `t2` has accessed). If `t2` commit first, the conflict is detected when `t1` commits and resolved by rolling back `t1`.

**Critical TVar**

A TVar is *critical* for a transaction `t1` if a modification to this TVars causes `t1` to roll back. In the current implementation, each time a transaction executes `readTVar`, the read TVar becomes critical for this transaction. Note that the time a TVar is critical for a transaction depends heavily on the implementation of STM.

**Deadlock**

The term *deadlock* is defined in different ways in literature. A deadlock is either a specific schedule or a property of the source code. In the course of this thesis a deadlock is a static code property. We assume that a thread is either `active`, `ready` or `suspended`. If the thread is `suspended`, it waits for an external signal to swap to the `ready` state. If it is ready the scheduler is allowed to swap this thread to `running` which grants the thread access to the CPU. If the thread executes specific operation (for example the semaphore operation `P`) it enters the state suspended. Otherwise it is up to the scheduler to swap the thread to the `ready` state and take away the access to the CPU. With kind of system, a deadlock is defined as follows: A program contains a deadlock if there is a schedule that all threads are suspended.

**External Semantics**

The *external* semantics of a program or function are the semantics that are **observable** from the hightest level of abstraction. Usually there is no need to distinguish internal and external semantics. This is only needed if the function is inspected on different levels of granularity. The external semantics always refers to the semantics relevant for the user. In this thesis we need this differentiation because we inspect the STM library on two levels of granularity. One is the user level. The other is the developer level, which we call internal semantics. For example the external semantics of `writeTVar`: *writes the passed value to the passed TVar*. The internal semantics on the other hand (in the original implementation) of `writeTVar`: *logs the passed value as newValue for the passed TVar*.

**STM Action and Transaction**

There is no difference between an STM action and a transaction. Every STM action can be executed with atomically and thus is a transaction. There may be exceptions, but I use the term transaction to denote an STM action that is supposed to be executed with `atomically`. I use the term STM action to speak about parts of transactions. For example:

```
withdraw src am = do
  balSrc <- readTVar src
  writeTVar src (balSrc - am)
```

`withdraw acc1 50` is a transaction, but `balSrc <- readTVar src` is an STM action. From the perspective of Haskells type system there is no differnce between these two parts. This is a convention not a definition.

**Code Conventions**

In the code examples givin in this thesis, I often use undeclared variables for the sake of space. The type of these variables can usually be derived from the context. Nonetheless, for clearity I will give a guideline for naming conventions. In many code examples `t1` and `t2` are used to denote TVars. The type of these TVars depends on the context, but is never important for the example. The examples that refer to the bank accounts (Section 2.1) use `acc1` and `acc2`. If `t1` or `t2` are used to denote transactions, this is always explicitly noted in the text. `a` and `b` denote pure values and are always introduced in the code.

All code examples use the do-notation, which is syntactic sugar for the monadic functions »= and ». In the text that explain the examples, we refer to »= instead of `<-`. In case you are not familiar with do-notation, it is highly recommended to take a look at the translation of do-notation to monadic operators, for example: `https://en.wikibooks.org/wiki/Haskell/do_notation`.

# Chapter 3

# Concept

In this Chapter we will explore an Approach to handle the problem of *Unnecessary Rollbacks* described in 2.3.1. Before we can understand a solution to this problem we need to analyze the technical reason more precisely. I was not able to find a satisfying solution for the problem of *Unnecessary Recomputations*. However, in Chapter 6 ideas to solve this problem are discussed.

## 3.1 Unnecessary Rollbacks

Remember the idea given in 2.3.1. We suggested to delay the `readTVar` operations to the commit phase rather than executing them immediately in the computation phase. While the idea works for the example of a normal `transfer`, the idea does not work for the following example:

```
limitedTransfer src dst am = do
  srcBal <- readTVar src
  if srcBal < am
    then return ()
    else do dstBal <- readTVar dst
            writeTVar src (srcBal - am)
            writeTVar dst (dstBal + am)
```

If we use this function, the result of `readTVar src` is needed in the computation phase and therefore the evaluation cannot be delayed to the commit phase. The value is needed to decide the condition of the `if` expression. To be exact the value is needed to determine the control flow.

This leads to the question whether there is a way to determine if the result of a `readTVar` effects the control flow or not. The current implementation does not do this. The main problem is the bind operator: `»= :: STM a -> (a -> STM b) -> STM b`. This operator allows us to extract the result of an STM action from the STM context, for example the result of a `readTVar`. This means the STM library loses any possibility to observe this value. The value is no longer in the libraries reach. Thus the library is not able to decide if the value is used to alter the control flow. Furthermore, the library is not able to determine if the control flow alters when the value is modified. The only way to guarantee the ACI properties is to restart the transaction when the TVar is modified. Otherwise the consistency may be violated.

If the library handles a value that is **not** used for branch conditions as if it were used for branch conditions, it may lose performance, but preserves the correctness. We already know an example for this. When we introduced STM in 2.1, we examined two transactions executing `transfer`. In 2.3.1 we detected that this would roll back at least one of these transactions.

If the library on the other hand handles a value that **is** used for branch conditions as if it were not, the library would not perform unnecessary rollbacks, but may violate the ACI properties. Thus GHC handles all values as if they are used for branch conditions to ensure the correctness of the implementation. Take for example `limitedTransfer acc1 acc2 50` and assume the initial value of `acc1` is 60. The transaction executes `read acc1` (and does not handle this TVar as a critical TVar) and determines the branch condition to be false. Then another transaction withdraws 20 from `acc1`. Afterwards the initial transaction resumes and completes its computation phase. Since no TVars are critical the validation succeeds and the transaction continues. The behaviour depends on the semantics of `readTVar`. Either the TVar is read now (in the commit phase) or the TVar is not evaluated, since it was already evaluated to determine the branch condition. If the TVar is read again and the update is commited, the new value of `acc1` is -10. Certainly nothing you expect, when executing a `limitedTransfer`. If the TVar is not read again the new value of `acc1` is 10. This is a lost update that generated money.

## 3.2 Approach

My approach to avoid unnecessary rollbacks and preserve the correctness is to handle all TVars uncritical, at first. While executing the computation phase, the TVars whose values are used to alter the control flow become critical. All `readTVar` operations are evaluated as late as possible, meaning a read on an uncritical TVar is executed in the commit phase and a read on a critical TVar is executed as soon as its value is used for some kind of branch, by which the TVar becomes critical.

Branch features in Haskell are the following:

- `if-then-else` expressions

- `case` expressions

- guards in functions or case expressions

- pattermatching in functions

Whenever a value is passed to one of these constructs, the associated TVars are marked critical and the affiliated read operations are evaluated. The *associated TVar* are the TVar that the value depends on. In other words, the TVars that must be read to determine the value. A value may depend on multiple TVars. Consider the following example:

```
transaction =
  a <- readTVar t1
  b <- readTVar t2
  if a < 5
    then if b - a < 0
            then ...
```

In this example the value in the condition of the first `if` expression depends on `a`, but not on `b`. The value in the condition of the second `if` on the other hand depends on `a` and `b`. If `a` is greater or equal to 5, only `t1` is critical for this transaction, otherwise `t1` and `t2` are critical for this trasaction. However, the time `t1` and `t2` are critical is different because `t1` becomes critical when the first `if` condition is evaluated and `t2` when the second `if` condition is evaluated. `a` is **not** evaluated again when the
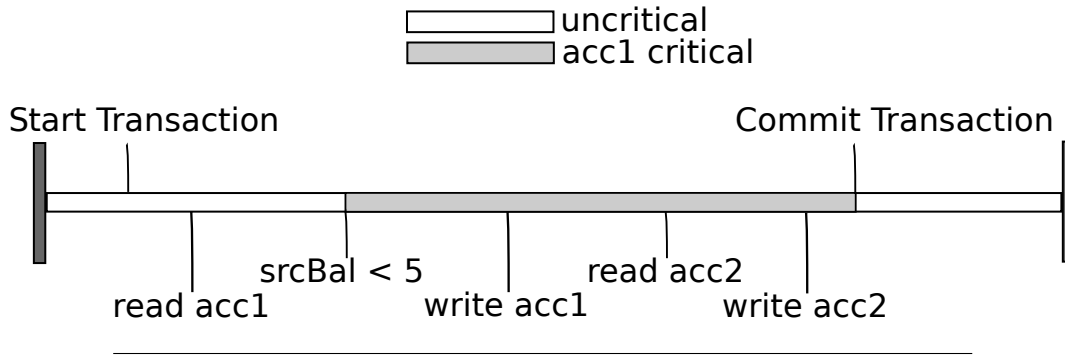
FIGURE 3.1: The critical time of the TVars in `limitedTransfer` with the alternative approach.

second `if` condition is evaluated. Every TVar is read at most once per transaction regarless the number of branches that depend on the TVar[1]. Even when the user evokes multiple `readTVar` operations on the same TVar, the actual TVar is accessed just once.

At the end of the computation phase all reads that are needed to decide the control flow are evaluated. All other reads are not relevant for the control flow and thus not evaluated. In the `transfer` example no `readTVar` is evaluated in the computation phase because neither of the TVars is used to decide the control flow.

Lets refrain from STM and concurrency for a second. This kind of evaluation is well known in Haskell. There are two cases where Haskell demands the evaluation of an expression[2]. The first is, if Haskell needs the value to execute a forreign function such as `putChar`. IO actions often contain calls to foreign functions. Although it is not necessary that an IO action calls a foreing function and evaluates its subexpressions, every IO action potentially evaluates its subexpressions. The second is, if the value is needed to decide a branch condition.

STM is an abstraction that allows us to use sequential programming in a concurrent context. So we can use the same evaluation strategy as in the sequential context. Even if the syntax suggests it, we do not explicitly specify when a TVar is read. We just specify the origin of the value and the evaluation is handled by the underlying system. Since the computation phase is processed in the STM monad, there are no IO actions allowed. This implies the only time we need to evaluate an expression is when we need to decide a branch condition. Everytime we execute other computations on TVar values and write them back or return them, this is not executed in the computation phase, because it is not needed. By evaluating only the reads that are needed and just before they are needed we minimize the time the TVars are critical. Figure 3.1 shows the effect on `limitedTransfer acc1 acc2 5`[3]. The TVar `acc2` is not critical at any point in the transaction, because it is not needed for the control flow. `acc1` on the other hand becomes critical at the time its value is used to evaluate the branch condition (`srcBal < am`).

When the commit phase starts some reads are evaluated and some are unevaluated. Before the remaining reads are evaluated the transaction acquires all locks for TVars it has modified. Chapter 4 explains that it is enough to lock the TVars that

---

[1]It should be clear that the actual TVars are read again, when the transaction is rolled back.

[2]There are other cases where Haskell demand the evaluation, but these are user defined strictness annotation suchs as `seq` or `!` for strict pattern matching or strict constructors. But these are not part of the actual semantics of Haskell.

[3]We assume that `srcBal` is greater than 5.

are modfied instead of all TVars that were accessed. Then the transaction is vali-
dated. This validation is similar to the validation in the original implementation; it
is a comparision between the values in the log (the TVars that the transaction needed
to read to determine the control flow) and the current values of these TVars. If all
values match, the transaction is valid, otherwise it is invalid. An invalid transaction
is rolled back immediately after the locks are released. If the transaction is valid, it
finally executes the remaining `readTVar` operations in order to modify the actual
TVars and determine the return value of the transaction. Then the modifications the
transaction performed are committed to actual TVars. By committing the modifica-
tions the old values in the TVar are overwritten. This is the reason the `readTVar`
operations can not be delayed any longer. Parallel to writing the actual TVars, the
locks for these TVars are released. The last step of the transaction is to return the
result.

# Chapter 4

# Implementation

In this Section we look upon the implementation of the aforementioned changes. Unlike to the original implementation this implementation does not use C primitives; it is a pure Haskell library. This brings some advantages and one disadvantages. The disadvantage is the performance as discussed in Chapter 5. For the costs of performance we gain a library that is easy to understand and extend. The original implementation is interwoven with the GHC runtime environment. Some STM functions are evoked by the scheduler to ensure the consistency. This makes the library sensitive to changes. To ensure the correctness of such a library is significantly harder than for pure library. In other words, the devlopment of a pure library is safer and faster. (Huch and Kupke, 2005) presented a high level Haskell implementation of STM. Their aim was to provide a pure Haskell implementation that is equivalent to original implementation of (Harris, Marlow, et al., 2005). Preceding this master thesis, I optimized that implementation. I replaced internally used data structures and performed two changes to the internal semantics. First, the initial high level implementation used a global lock to synchronize concurrent transactions. This coarse grained locking was substituted by a fine grained locking. Instead of a single global lock, each TVar holds its own lock and transactions acquire the minimum amount of locks to commit. This allows transactions that do not conflict to commit simultaneously. Second, I altered the conflict detection. The initial implementation used a validation process similar to the GHC implementation. In the optimized implementation, each TVar has a queue associated. If a transaction reads a TVar, it enters a reference to itself to this queue. If a transaction successfully commits a change to a TVar, it notifies all transactions in the associated queue. If a transaction is notified, it is rolled back. This way of conflict detection has the advantage that conflicts are detected earlier than before. This implementation is called the *Project implementation* in the following. Figure 4.1 visualizes the development process of the libraries. In the following Section we give a detailed description of the implementation developed in the course of this thesis.

## 4.1 STM Types

Before we head over to the implementation of the external interface of STM, we investigate the types that are used in this implementation; starting with the STM type itself:

```
data STM a = STM (StmState -> IO (STMResult a))
```

In its core the STM data type is similar to a state monad. The IO type was initially needed to perform the reads in the computation phase. As we will see in the next section, in this implementation it is needed only to create new TVars in the computation phase. `readTVar` and `writeTVar` do not need to process IO operations. Thus
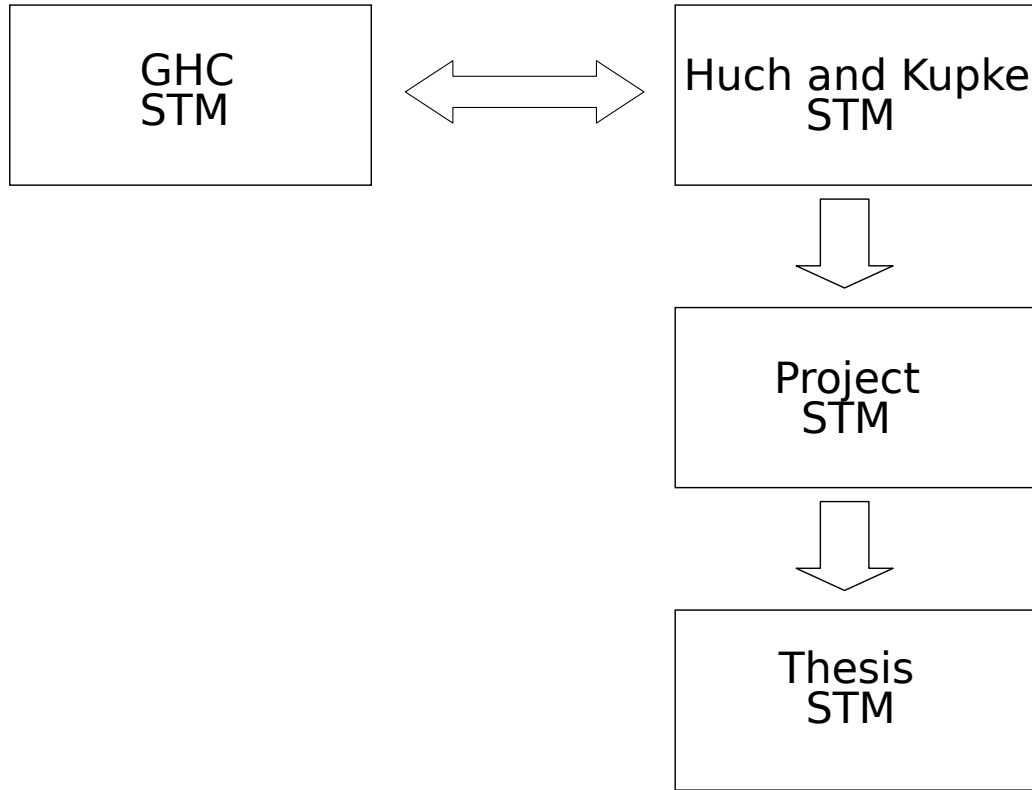
FIGURE 4.1: The STM implementations for Haskell.

an STM action takes a state and returns a result depending on this state. There are three possible results:

```
data StmResult a = Retry StmState
                 | InValid
                 | Success StmState (Maybe a)
```

The first constructor is used to indicate that `retry` occurred. This must be distinguished from `InValid`, since `orElse` as well as `atomically` react differently on these results. If `Retry` is returned, `atomically` validates and rolls back at an appropriate time; `orElse` starts the second transaction. This is the reason `Retry` is accompanied by the state. The state holds the necessary information about the TVars that the transaction has read. This allows to validate and possibly suspend on particular TVars. `InValid` on the otherhand indicates always that the transaction is not valid and thus must be restarted. The last constructor is the desired outcome of an action. If the transaction does not fail, it returns `Success` and a state as well as the result. As before the state is needed for validation. The result is wrapped with the `Maybe` type. The value `Nothing` never occurs. The reason for this wrapping is explained in Section 4.2.

The core data structure of STM in Haskell is the `TVar`:

```
data TVar a = TVar (MVar (IORef a))   ――value
                   ID
                   (MVar [MVar ()])   ――queue
                   (MVar ())          ――lock
```

All TVars have a globally unique identifier called `ID`, which is immutable. The `MVar (IORef a)` holds the actual value of the TVar. The MVar is used to allow multiple

threads to access and modify the same data structure. This is not possible with pure values in Haskell. The inner `IORef` is needed to enable a correct validation, which is explained in detail in Section 4.2. `MVar [MVar ()]` is the queue of the MVar that is needed to observe this TVar. Remember that this is important to delay the rollback when `retry` is evoked. The last part is an explicit lock for this TVar. Currently there are two implementations as a result of this thesis. One of these implementations locks a TVar via the explicit lock and the other by taking the MVar that holds the value. The details are explained in the next section.

The computation phase does not allow any observable side effect and its only result is the `StmState` and a single value. Hence, the core data structure is the `StmState` which holds all information that are needed to commit a transaction:

```
data StmState = TST { writeSet   :: IntMap ( Maybe (),
                                             Maybe (),
                                             MVar (IORef ()),
                                             MVar () ),
                      notifies    :: IO (),
                      readSet     :: ReadSet,
                      retryMVar   :: MVar (),
                      uEReads     :: [ Maybe () ]}
```

The `writeSet` is similar to the log in the GHC implementation. Here it is an `IntMap` [1]. The keys are the IDs of the associated TVars. The elements contain the `expectedValue`, the `newValue`, the `actualMVar`, and the explicit lock. The `actualMVar` is the value holding MVar and is needed in the commit phase when the transaction sucessfully commits to publish its modifications. The `expectedValue` and `newValue` are similar to the log entries presented in 2.2, but their purpose is slighty different. The `Maybe` type indicates that the values also can be `Nothing`. This holds only for the second entry. The first entry has the `Maybe` type for the same reason `Success` has it. The second entry on the other hand can become `Nothing`. This indicates that the TVar was read but not written by the transaction. Note that there are two kinds of *read*. The first is that a `readTVar` operation is porcessed in the computation phase and the second that the current value of the TVar is read. In the original implementation these two occur always at the same time. In the new approach theses kinds of reads occur at different times. For clearity, from here on I will use IO-reads when I refer to a read from the actual IORef to access the current value of that TVar. Since Haskell does not provide a simple solution for inhomogeneous containers, I decided to use the type `()` and cast all values via `unsafeCoerce` before entering them in the `writeSet`. Whenever a value is taken from the `writeSet` it is casted back to its original type by `unsafeCoerce`. The ID of the TVar ensures that the type it was casted from does not differ from the type it is casted to.

`notifies` holds an IO action that is process when the transaction successfully commits. This IO action notifies all transactions that are waiting on a TVar that is modified by the committing transaction.

`readSet` stores information about the IO-reads that were performed. The type `ReadSet` is defined as follows[2]:

```
type ReadSet = IORef ( InMap ( IORef (),
                               MVar (IORef ()),
```

---

[1]This refers to the IntMap in the standart libraries of Haskell: `https://hackage.haskell.org/package/containers-0.5.8.1/docs/Data-IntMap-Strict.html`

[2]Notice the difference between `readSet` and `ReadSet`. `readSet` is the field of the `StmState` and `ReadSet` is the data type.

MVar [MVar ()]))

The need for the outermost IORef is explained later. This type is like the `writeSet` an `IntMap`. The keys are the IDs of the TVars. The entries consist of three parts. First, the value (in form of the IORef) that was present when the IO-read was executed. Second, the MVar that holds the value of the associated TVar. Third, the queue of the associated TVar. The exact ussage of this is explained in the next section.

`retryMVar` is an unique MVar for every transaction. This is the MVar that is entered into the queues of the TVars when the transactions processes `retry`.

`uEReads`(unevaluated read) contains all unevaluated IO-read operations. This is essential to be able to process the IO-reads that where not evaluated in the computation phase. The `Maybe` type is again used for the same reasons as it is in `StmResult`.

This concludes the overview on the data structures used and we now have the necessary foundation to understand the implementation of the interface Functions.

## 4.2 Interface Functions

**newTVar**

This is the only function (besides `atomically`) that uses IO actions. A TVar consists of three MVars and one ID. To ensure the uniqueness of the IDs the STM library defines `globalCount :: MVar Int` and a function `getGlobalId :: IO Int` that takes the `globalCount`, increments it, writes it back and returns the old value. Due to the semantics of the MVar only one thread at a time can get an ID. No ID is assigned twice unless `globalCount` overflows. Because the chances that this happens are narrow, no countermeasures are performed to detect or avoid this problem.

All MVars are created with `newMVar`. This means all MVars are initialized when the TVar is created. The queue holds and empty list the lock holds unit and the value holds an IORef with the passed value. In the end the unchanged state as well as the new TVar is returned.

Lets take a look at the following example:

```
transaction = do
  tv <- newTVar 5
  a <- readTVar tv
  writeTVar tv (a + 1)
```

The first action of this transaction is the creation of a new TVar. This means a new IORef with the value 5 is created. After that three new `MVars` are created. The first contains the IORef, the second `[]` and the third `()`. With the value gained from `getGlobalId` a new TVar is created. We assume the ID is 1 for now. The state on the other hand is not modified. Then the second action, `readTVar tv`, is processed.

**readTVar**

This function is the core of the new implementation. One part of `readTVar` is similar to the original Implementation. If the TVar is present in the `writeSet`, the `expectedValue` is returned (explained together with `writeTVar`). The differences are in the other part of `readTVar`. Recall the type of `readTVar :: TVar a -> STM a`. The result is an STM action. This means a function that takes an `StmState` and returns an `StmResult`. Unless the result is `InValid`, it contains an `StmState`. In other words, `readTVar` transforms the state and returns the value of the TVar. To return the value of the TVar requires to IO-read the TVar. To avoid the IO-read at

this point, it is not executed directly. The function `buildVal` helps us to delay the evaluation.

`buildVal` wraps the value with `unsafePerformIO`. `unsafePerformIO` allows us to execute IO actions in pure Haskell code. These IO actions are executed when the evaluation of the expression that they occur in is demanded. In this case the IO action consists of two parts. First, it IO-reads the current value of the TVar and returns it after wrapping in into `Just`. The wrapping with `Just` is explained when we examine the implementation of atomically. Second, entering the information that are needed to validate to the `readSet`. These are the ID of the TVar as key, the MVar of the TVar that holds the value, the value that the transaction saw when it IO-reads the TVar, and the queue of that TVar. The value the transaction has seen is created by the first part of the `unsafePerformIO` action. The other information are the arguments of `buildVal`. After `buildVal` created this wrapping, the expression is (from Haskells view) a normal expression, but the IO-read is performed when the evaluation of this expression is demanded and not earlier; simultaneously is this evaluation logged in the `readSet`. Haskell demands the evaluation of values just to decide on branch condition or to perform foreign functions. IO actions that demand foreign functions are not allowed in the STM monad by the type system. Thus the evaluation in the computation phase is only demanded when the control flow depends on it.

Back to `readTVar`. `buildVal` is used to create the result. `readTVar` also modifies the the `StmState`. It adds the newly created value to `uEReads`. And it extends the writeSet. At this point there is no entry for the TVar in the writeSet yet, otherwise this part of `readTVar` would not be executed. A new entry is inserted. This entry contains the constructed value and the value holding MVar of that TVar. The `newValue` field of this entry is `Nothing`. In the end the constructed value and the new state are returned in the `Success` constructor.

Reconsider the example givin in the previous section:

```
transaction = do
  tv <- newTVar 5
  a <- readTVar tv
  writeTVar tv (a + 1)
```

We explained already the first action. The first step of the second action is to lookup `tv` in the `writeSet`. Since `tv` is not present in the `writeSet`, `buildVal` is called to create the `unsafePerformIO` action. The result (called `val` in the following) is added to the `writeSet`. Thus the `writeSet` is no longer empty. It contains an entry for `tv`:

```
writeSet1 = {(1,(val, Nothing, mvar))}
```

`mvar` is the MVar of `tv` that holds the IORef with the value. The `1` is the ID of `tv`. Aditionally is `val` added to `uEReads`. All other parts of the state are unchanged. Due to `<-`, `val` is bound to `a` in the end of the second action.

**writeTVar**

This implementation is straight forward. Two modifications to the state are performed. One is the modification of `notifies`. A successful commit of the transaction means a modification to the TVar that `writeTVar` was called on. Thus this transaction possibly needs to notify waiting transactions. An IO action is created, that notifies all transactions in the queue of the TVar. This IO action is sequenced

by » with `notifies` of the initial `StmState`. The resulting action then notifies all
TVars that are written by this transaction. This may lead to an action that notifies the
same queue twice. This is a minor issue because after the first notification the queue
is emptied. To avoid this, you need to lookup the value in the writeSet, to see if it
was already written and if so do not extend the `notifies` action. This look up is
most likely more expensive than a notification on an empty queue (this has not been
investigated).

The second modification that `writeTVar` performs is the modification of the
`writeSet`. After wrapping the value for type reasons with `Just`, it is entered in the
`expectedValue` and `newValue` field of the associated entry in the `writeSet`. The
last field of the entry is the MVar of the TVar that holds the value. There is a reason
to set both, the first and second field, to the new value. If the second field is **not**
`Nothing`, it indicates that this value is modified by the transaction and in the case
the transaction sucessfully commits, the value in this field is written. The first field
of the entry is the field that `readTVar` returns if it finds this entry. To avoid this
doubled entry, the obvious choice were to use pattern matching to check whether
the second value is `Nothing` and if not return that value. This does not work in this
implementation. The problem is that the pattern matching forces the evaluation. By
forcing the evaluation of an expression you may evaluate IO-reads and thus make
TVars critical (IO-read them) that are not critical. This does not happen in the current
implementation because in every entry of the `writeSet` the first field is always not
`Nothing` and thus can be returned without pattern matching on this value. In the
end of `writeTVar`, the new state and `()` are returned.

For the example it means that the `writeSet` and `notifies` are modified by the
third action. The `writeSet` looks like this after the third action[3]:

```
writeSet2 = {(1,(val + 1, val + 1, mvar))}
```

In both, the first and the second field, `(val + 1)` is entered. `notifies` is initially
`return ()`. After the `writeTVar` action it is:

```
notifies1 = do
    queue <- takeMVar waitQ
    mapM_ (flip tryPutMVar ()) queue
    putMVar []
    return ()
```

`waitQ` is the queue of `tv`. The notification is performed by trying to put `()` to
the MVars in the queue. It is important to use `tryPutMVar` instead of `putMVar`.
Otherwise the implementation contains a high potential for deadlocks (imagine two
transactions try to notify the same transaction at the same time). In the end the queue
is emptied to avoid another transaction to notify the same transactions again. The
final `return ()` is the initial value of `notifies`. The result of the three action in
the example is:

```
Success state ()
```

`state` is a `StmState` that contains `notifies1`, `writeSet2` and the modified
`uEReads`. The `readSet` and `retryMVar` are similar to the one in the initial state.
This state is used in atomically to commit the changes.

---

[3]In fact `val` is of type `Maybe Int` and thus cannot be added to `1`. The bind operator on `STM`
unwraps the value, so we can use it as pure value. This means in the case, that the actual entries are
`Just (fromJust val + 1)`. For the sake of comprehensibility, we ignore this for now.

**atomically**

This is the heart of every STM implementation in Haskell; it allows the user to execute STM actions in a transactional manner. Before `atomically` starts the computation phase, it creates an initial state. This state basically holds no information and a commit of this state results in no change at all. After the initial state is created, the computation phase starts. In this phase atomically does not do anything but pasing the initial state to the STM action.

After the `StmResult` is calculated, atomically starts the commit phase by interpreting the result. There are three possible outcomes. First, the result is `InValid`. If this occurs the computation phase is just restarted with the initial state[4].

Second, the result is `Retry newState`. This causes the transaction to lock the TVars in its `readSet` which can be accessed via `newState`. After the locks have been acquired, the `readSet` is validated. If it is not valid, the locks for the TVars are released and the transaction is rolled back. If the transaction is valid, the transaction enters its `retryMVar` to the queues of the TVars it has read. These queues are stored in the `readSet`. Then the transaction realeases the locks and executes a `takeMVar` on its `retryMVar`. This lets the transaction suspend until another transaction notifies it. After the transaction was notified, it removes its `retryMVar` from the queues (if that has not already happened) and restarts.

Third, the result is `Success newState result`. This is the most interessting case because it is the commit phase and thus leads not necessarily to a rollback, in contrast to the other two results. There are currently two different implementation as a result of this thesis which only differ in this part. The performance tests presented in Chapter 5 provide a clear result on which implementation is better in general. Nonetheless, I will present both implementation. If it is not clear which implementation we are talking about, we will refer to the first implementation as STMLA (STM lock all) and the second STMWSL (STM write set lock).

STMLA starts the commit phase by locking all TVars is has accessed. This implementation uses the explicit locks of the TVars. This information about which TVars were accessed are stored in the `writeSet` since `writeTVar` as well as `readTVar` extend the `writeSet`. When all locks are acquired, the transaction validates its `readSet`. If this is not valid, the transaction is rolled back. If it is valid, the transaction evaluates all expressions in `uEReads`. This is done by the `seq` function which forces the evaluation of its first argument and returns its second argument. As always in Haskell the expression is evaluated to WHNF[5]. This is where the `Maybe` type comes in handy. `buildVal` created a `Maybe` value that was wrapped with an `unsafePerformIO` action. If we force the evaluation of this expression, the IO-read is performed, but the underlying expression is not evaluated. It is wrapped with the `Just` constructor. Without this constructor, we would evaluate the expression that is stored in the TVar. Imagine this expression is a complex computation. This means, we suffer a performance overhead because the computation is not yet needed (or not needed at all). Even worse, we execute this in the commit phase while we hold the locks for many TVars. To maximize parallelism, we aim to minimize the time the TVars are locked. Besides these performance problems, there is a serious semantic problem. Everytime we force the evaluation of an expression that is not needed, we change the semantics of the program. This may lead to exceptions or loops that

---

[4]Actually the `readSet` needs to be discarded explicitly, because it is an IORef and thus would not be empty by taking the initial state.

[5]In case you do not know what *weak head normal form* is please refer to `https://wiki.haskell.org/Weak_head_normal_form`

would not occur normaly. After the `uEReads` are evaluated, the actual writes are prepared. The writes are divided in two parts. First take the value from the MVar and second write the new value. Preparing means all the value holding MVars of all TVars that are going to be written are emptied. Then the `notifies` are processed. At last the new values are written to the value holding MVars. This dividing of the writes is needed to prevent a very subtle bug. The last step of the commit phase before returning the result is to realease the locks that were taken at the start of the commit phase.

The second implementation (STMWSL) starts the commit phase by validating. As before, the transaction is rolled back if it is invalid. If it is valid, the remaining reads in `uERead` are evaluated. After all reads have been processed, the actual writes are prepared. This means the TVars that are modified by the transaction are accessed and their value holding MVar is emptied. At this point no other transaction is able to either read or write these TVars. In the first implementation the explicit locks were taken. This does not prevent other transactions from reading these TVars. After the writes are prepared, the transaction is validated again. In contrast to the first validation, this validation is mandatory. It is possible that the current values of the TVars have changed after the reads were processed and before the writes were prepared (and by this the TVars locked). The first validation is added to avoid the problem that the transaction acquires the locks, although it is invalid. The costs for locking are much higher than the costs for validation. If the transaction is invalid after the reads are processed and the transaction acquired the locks, it is rolled back. The second validation ensures the correctness and is similar to the one in the first implementation (STMLA). After the second validation, the `notifies` are processed. Then the writes are processed. At last the result is returned. Note that it is not necessary to unlock the TVars explicitly because only the value holding MVars of the TVars serve as a lock in STMWSL. By filling these MVars the associated TVars are unlocked.

Both implementation have their own benefits. STMLA does never roll back when the transaction does not branch at all. If two transactions read the same TVar they cannot commit at the same time in STMLA. In STMWSL it is possible that two transaction, that read the same TVar, commit simultaneously (if non of them also writes this TVar). It is also possible that a rollback occurs, even if the transaction does not branch. The problem is that the remaining reads are evaluated in an unlocked context. This means that other transactions can modify these values after they were read and thus invalidate the reading transaction. It would be better, if the transaction process the `uEReads` after it has locked the TVars. This is currently not possible as we will discuss in detail in the next section. When we only lock the modified TVars, it is essential that no other transaction reads them while we hold the locks. Otherwise the ACI properties cannot be guaranteed. If we take the value holding MVars it prevents other transactions from reading the associated TVars. Unfortunately, this also prevents us from reading the TVars. That is why we need to process `uEReads` before we acquire the locks for the TVars. The clever reader may have noticed that we get the values of the TVars when we prepare the writes because we execute `takeMVar` on the value holding MVars. To use these information would require the `unsafePerformIO` actions produced by `buildVal` to behave diffrently when it is executed in the computation phase and in the commit phase. Nevertheless, this topic remains for future work and is discussed in Chapter 6.

A solution that combines STMLA and STMWSL would be better[6]. A combined solution locks the TVars that the transaction has modified in a manner that no **other**

---

[6]This could not be proved because no such implementation is currently available.

transaction can read them. The transaction itself is still able to read the locked TVars. Then the transaction is validated. If it is valid, the remaining reads are processed and the transaction publishes its modifications. At last the transaction returns the result, after unlocking the TVars.

If we execute `transaction` introduced in Section 4.2 with `atomically` (of STMLA), the computation phase produces the previously mentioned state. The `writeSet` in this state contains a single entry. The entry for the newly created TVar. The commit phase starts by validating the `readSet`. Since the `readSet` is empty, the validation succeeds and the transaction acquires the lock for `tv`. After the lock is acquired, the transaction evaluates `uEReads`. `uEReads` contains one value, namely `val`, the `unsafePerformIO` action that reads `tv`. The evaluation of this value has (thanks to sharing) an effect on the `writeSet` as well:

```
writeSet = {(1,(Just (5 + 1), Just (5 + 1), mvar))}
```

The next step is to prepare the writes. This means `mvar` is emptied via `takeMVar`. Before it is filled again, `notifies` is executed. `notifies` does nothing because the `waitQ` of `tv` is empty. Finally the expression `(5 + 1)` written to `mvar`. Note that the expression itself is not evaluated. The last step before returning `()` is to release the lock for `tv`.

### »=

The monadic bind operator, »=, is an important part of the interface because it allow us to use the do-natation. The aim of this thesis is to provide an alternative STM implementation for Haskell without altering the semantics or interface. One of the most comfortable things about STM is the way we use it. Thus it is important to retain this feature. `»= :: STM a -> (a -> STM b) -> STM b` allows us to extract the result of an `STM` action from the STM context and use it to create a new STM action. Additionally it is used to translate `<-` of the do-notation. The implementation is straight forward, but for the sake of completeness presented nonetheless. The result is a STM action, meaning that it is a function that takes a `StmState` and returns an `StmResult`. Thus the function »= basically has three arguments. The `STM` action, the function of type `a -> STM b`, and the `StmState`. These arguments are called *passed action/function/state* in the following.

The resulting function first applies the passed state to the passed action. If the result is `InValid`, it just returns `InValid`. If it is `Retry newState` it returns `Retry newState`.

Note that it is important to pass the `newState` created by the action because it contains information that are needed to wait before rolling back.

If the result of the action is `Success newState res`, the function first needs to unwrap `res` (it is wrapped with the `Just` constructor). This is done by the partial function `fromJust`. The implementation ensure that `Nothing` can never occur and thus the call to `fromJust` is safe. After the value is extracted, it is applied to the passed function. This application results in an `STM` action. To gain the `StmResult` as the result of the function the `newState` is applied to the action. Thanks to the non-strict semantics of Haskell the call to `fromJust` is not evaluated immediately. This is the key of this implementation, since we avoid the evaluation of the IO-read by this. If `fromJust` would be evaluated when the bind operator is evaluated, a action like `a <- readTVar t1` would not differ from the original implementation because the `unsafePerformIO` action would be executed immediately.

**retry and orElse**

The implementation of `retry` and `orElse` are very simple. `retry` is a STM action and thus gets a state passed. The passed state is wrapped with the `Retry` constructor. That is all `retry` does.

   `orElse` executes the first action with the passed state and if the result is not `Retry newState` it just returns this result. If it is `Retry newState`, the second action is executed with the passed state (not the `newState`) to discard the writes executed by the first action. Since the `readSet` is an `IORef`, we do not lose the information on the TVars that were read.

## 4.3   Notes on the Implementation

Some details of the implementation remained unexplained until here. These details are highlighted in this section.

**Deadlock Avoidance**

One Motivation for STM is the deadlock freedom. We have seen in Section 2.1 that acquiring multiple locks always includes the danger of deadlocks if multiple threads work concurrently. There are two concepts to avoid deadlocks for such settings. The setting is that multiple threads try to acquire and arbitrary number of shared locks. The first avoidance strategy is, if thread tries to acquire a lock that is not available, the thread releases all locks and tries to acquire all locks again. This is continued until the thread can acquire all desired locks. The other avoidance strategy is to acquire the locks in a global order. STMLA as well as STMWSL use the latter. The TVars that need to be locked are stored within `IntMaps`. If a transaction trys to lock these TVars, the transaction converts the `IntMap` to a list and processes it. Since the conversion from an `IntMap` constructs an ordered list, the transactions always lock in a global order. The TVars are ordered by their IDs.

### 4.3.1   STMWSL

The double validation of STMWSL seems to be an unnecessary overhead. However, the original implementation also uses this scheme which was proposed by Fraser (Fraser, 2004, Page 42). The high costs for locking the TVars are the reason. This locking itself is not particular expensive, but it hinders parallelism. Everytime a TVar is locked no other transaction is able to read, lock or validate this TVar. Additionally if a transaction tries to access a locked TVar, it is suspended and a context switch follows. Context switches in Haskell are not as expensive as context switches of OS threads, but they schould not be neglected, especially when dealing with a large number of threads. Validation is an operation which needs two memory accesses per entry in the log. One memory access is the access to the log entry to look up the expected value and the other memory access is the access to the actual TVar[7]. This is reasonable considering that the log only consists of TVars that are needed to determine the control flow. So validation is significantly faster than locking.

---

[7]The log entry needs to be accessed two times. First to get the expected value and second to get the next entry because the log is a dynamic structure and thus needs a pointer to the next entry comparable to a list. Nevertheless the entry itself is (most likely) in one block in the memory. This is why I consider one memory access to be enough for the entry.

**Nested TVar**

To use newly created TVars it is necessary to share these with other threads. The mean to exchange data in STM are TVars. Hence it is not uncommon to have nested TVars when using STM. When using nested TVars in the alternative implementation it is important to know which TVars are evaluated during the computation phase. Take a look at the following example:

```
main = do
  tv <- atomically $ newTVar 42 >>= newTVar
  atomically $ readTVar tv >>= readTVar
```

This example creates a nested TVar in one transaction and read the inner TVar in the second transaction. In the second transaction the outer TVar is critical because the second `readTVar` uses pattern matching to extract values from the inner TVar. Thus the value of the outer TVar is used within a branch condition. Although the pattern matching of `readTVar` has only one case Haskell forces the evaluation of the the the expression that is passed to `readTVar`. The value of the inner TVar on the other hand is not used and thus not evaluated. Hence, the inner TVar is not critical.

This can be extended for further nested TVars as in the following example:

```
main1 = do
  tv <- atomically $ newTVar 42 >>= newTVar >>= newTVar
  atomically $ readTVar tv >>= readTVar >>= readTVar
```

In this example the outermost and the second outermost are critical in the second transaction. In general all but the innermost TVar are critical. Furthermore, this holds for all kinds of unfailable pattern matching. These force the evaluation of the passed expression although unfailable pattern matching is not a branch. Thus the following example also evaluated the IO-read and makes the TVars critical:

```
main2 = do
  tv <- atomically $ newTVar (4,2)
  atomically $ do
    p <- readTVar tv
    if one p
       then return 3
       else return 5

one :: (a,b) -> Bool
one (_,_) = True
```

We first create a TVar that hold a pair. In the second transaction we pattern match on this pair with the function `one`. The value of the pair are not used in `one` and the pattern matching cannot fail. Nevertheless, the evaluation of `p` is forced. Thus `tv` is critical in the second transaction. Note that these semantics have nothing to do with my implementation. This is an implementation issue of the lazy evaluation in GHC. To explain the evaluation strategy of GHC in detail is not aim of this thesis. To use STM it is important to note that even unfailable patter matching forces the evaluation of its argument.

# Chapter 5

# Evaluation

In this chapter we elaborate the impact of the implemented changes on the performance. Evaluating STM is worth a thesis itself. The biggest advantage of STM, the usability, is its biggest disadvantage with regards to performance testing. STM is a universal tool. Most synchronization problems can be solved with STM. This on the other hand means there is no clear way to test STM, especially when measuring the performance. Thanks to the moderate interface, testing the correct behaviour of STM is practicable. Although testing is no guarantee that the implementation is correct in all regards, but it narrows the space for bugs. I wrote a number of small tests to test the implementations for specific bugs. However, we will not investigate these correctness tests.

Unfortunately, testing the performance is extremely difficult. There are unlimited possibilities to use STM. Thus it is not possible to test the performance in general. We can only test the performance for specific cases. Thus it is impossible to say which implementation has the best performance in general. To be able to classify the differnt tests and to compare the differnt test to each other, we use the following properties:

- costs of a transaction

- level of currency

The *costs* of a transaction denotes the time the transaction needs to execute when no other transaction is present. It is important to distinguish this from the time the transaction needs to execute. The time heavily depends on other transactions. When a transaction is executed in a system with many other transaction that work on the same TVars, the chance that it is rolled back is higher than if the transaction is the only transaction in the system. In our case we meassure the cost of a transaction by counting the `readTVar` and `writeTVar` operations. The costs of a transaction do not depend on the level of concurrency.

The *level of concurrency* denotes the density of the TVar usage. A high level of concurrency is givin if there are many transactions which are reading and writing the same TVars. Read-only TVars are not considered, because reading them cannot result in a rollback. Furthermore, if every transaction works on different TVars, there is no concurrency at all. If there is only one transaction, there is also no concurrency. Thus only if all three requirements are satisfied, we speak of a high level of concurrency. The higher the level of concurrency the higher the chance a transaction is rolled back in the original implementation. This cannot be meassured in numbers, but we use this term to compare different tests in a relative way.

These properties are not statically measurable in general because they often depend on the state of the TVars. The cost may vary depending on branches. The level of concurrency may also depend on branch conditions because it determines

which TVars are accessed. Furthermore, the scheduler does affect the level of concurrency. The GHC runtime system uses the *round robin* scheduling scheme. If all transactions are sufficient cheap, they finish before their time expires. This means irrespective of the number of threads, there is no concurrency at all (if a single OS thread is used). To avoid these problems, the tests do not use branch contditions at all. Hence STMLA never rolls back.

## 5.1   Test Setup

Before we head over to the results, we will look upon the tests that are used to measure the performance. I used two tests to compare the different implementations. The frist test is called *StmTest*. It is used to test the performance and the correctness at the same time. StmTest has four parameters to control the costs of a transaction and the level of concurrency:

- `threads`

- `iterations`

- `tvars`

- `changes`

`threads` determines the number of threads that are working parallel. `iterations` is the number of transactions each thread executes. `tvars` is the number of TVars that are created and used. `changes` is the number of write and read operations per transaction. We use these parameters in the following as variables for integers. StmTest starts by creating `tvars` TVars. Then `threads` threads are created. Each thread chooses `changes` TVars randomly from all created TVars (the same TVar may be chosen multiple times). Then the thread starts a transaction. In this transaction the thread reads each of the chosen TVars, increments their value and writes the new value back to the TVar. This is repeated `iterations` times. When every forked thread has finished the main thread reads all TVars and sums their values. If the STM system is correct, the sum is (`threads * iterations * changes`).

By altering the parameters the level of concurrency and the costs per transaction can be controlled. More `threads` and less `tvars` result in a higher level of concurrency. More `changes` mean not only higher costs of a transaction, but also a higher level of concurrency. Unfortunately, in this test we are not able to increase the costs per transaction without increasing the level of concurrency. The overall runtime of the test can be managed with `iterations`. It also helps to avoid that the thread finishes before its round robin time expires. This test clarifies the previously mentioned problem of testing STM. These parameters can arbitrarily chosen and all of configurations are correct uses of STM. The number of configurations on the other hand is nearly unlimited. Thus it is not possible to compare the implementations with all possible configurations. To determine which STM implementation is the best overall is anything but trivial. Nevertheless, we use this test to compare the implementations on specific configurations to see their individuel strengths and weaknesses. Note that the transactions in this test always write all the TVars they have read. This is not necessarily the case when STM is used in real programs. That is why we use a second Test to meassure the performance.

*PerformanceTest* is a test to meassure the performance and not the correctness. In contrast to StmTest is has not four but five parameters:

- threads

- iterations

- tvars

- rWRatio

- writes

The first three parameters are the same as in StmTest. `rWRatio` determines the ratio between reads and writes. For example, if the `rWRation` is 5, the threads perform five reads for each write. `writes` on the other hands specifies the number of writes that are executed in each transaction. This test allows us to increase the costs of the transactions by increasing the level of concurrency only slightly. If multiple transactions read the same TVar there is no conflict. If multiple transaction on the other hand read and write the same TVar, there is a conflict. Hence, by increasing `rWRatio` the level of concurrency increases slightly while the costs for the transactions increase moderately.

PerformanceTest ist similar to StmTest. It first creates `tvars` TVars. Then it forks `threads` threads. Each of these threads creates `writes` lists ramdomly with `rWRatio` TVars in each. For each list the transaction reads all TVars, sums their values and writes them back to the first entry of that list. All lists are processed in a single transaction. This procedure is repeated `iterations` times. Since the TVars are choosen randomly, we cannot determine if the final state of the TVars is correct after all threads have finished. Another difference between the two tests is that StmTest uses a list to store and lookup the TVars while PerformanceTest uses an `IntMap`.

To meassure the execution time of the implementations the unix `time` command is used[1]. The test (either StmTest or PerformanceTest) are compiled with GHC 8.0.1[2] and the compiler flags `O2` for optimizations and `-threaded` to allow threaded runtime. The tests are exectued with the runtime option `-N` to allow multiple (in this case four) OS threads. The tests are executed on a system with an Intel(R) Core(TM) i7-6500U CPU @ 2.50GHz, 8GB @ 1600 Mhz DDR3 and a Fedora 25 OS.

## 5.2 Results

Finally we inspect the results of the performance tests. I performed two series of tests to compare the different implementations. In the first series the tests are configured by hand and the level of concurrency as well as the costs per transaction are modified. The total workload was the same in all tests in this series. In the second test series the number of threads are increased in every test. The workload per thread is unchanged and thus the total workload increases with the number of threads.

### 5.2.1 First Test Series

The results of the first test series are shown in Figure 5.1. The first column contains the test and its configuration. StmTest(threads,iterations,tvars,changes) means the `StmTest` was applied with the previously explained configuration parameters. PerTest(threads,iterations,tvars,rWRatio,writes) is the same for `PerformanceTest`.

---

[1]`https://en.wikipedia.org/wiki/Time_(Unix)`
[2]`https://www.haskell.org/ghc/download_ghc_8_0_1`

|                          | GHC    | Project | STMLA  | STMWSL |
|--------------------------|--------|---------|--------|--------|
| StmTest(20,1000,200,050) | 3.1780 | 3.5645  | 3.5340 | 3.6655 |
| StmTest(20,2000,200,025) | 3.3845 | 3.5700  | 3.6335 | 3.6665 |
| StmTest(20,0500,200,100) | 3.2978 | 3.8540  | 3.5905 | 3.7910 |
| PerTest(20,500,200,10,10)| 3.0420 | 3.2830  | 2.9225 | 3.4920 |
| PerTest(20,500,200,20,05)| 3.0670 | 3.3110  | 2.9425 | 3.4445 |
| PerTest(20,500,200,05,20)| 3.1520 | 3.4335  | 2.9455 | 3.3500 |

FIGURE 5.1: Average runtime in seconds for tests with a total work-
load of 1.000.000 `readTVar` operations.

The first row introduces the different implementations. GHC is the current library in GHC 8. Project is the highlevel library that this thesis is based on. STMLA and STMWSL are described in the previous chapter. For unknown reasons STMWSL performs poorly in this tests.[3]

In the first three tests we examined how changes in the costs of a transaction and the level of concurrency effect the runtime of the four systems. The number of threads and TVars are fixed. By altering the `changes` parameter, we increase the workload per transaction. To preserve the total workload, we decrease the number of iterations each time we incease the number of changes per transaction. In the first test we see that STMLA and the Project implementation perform equally good while the GHC library performs slightly better and STMWSL slightly worse. In the second test the workload per transaction is halved compared to the first test. The runtime of the GHC library and STMLA raises. Nevertheless, the GHC implementation is still the fastest. The increase in runtime can be explained with the fix overhead of each transaction. The second test executes twice as many transactions as the first test, thus it needs to start and commit twice as many transactions. The third test examines the other direction. The workload per transaction is doubled compared to the first test. This leads to an significant increase in runtime for the Project library and STMWSL. The GHC library becomes slightly slower with this configuration and STMLA remains on an equal level. This is the result we expected, since the rollback of a transaction is more expensive than in the first test. Fortunately, STMLA does not perform any rollback in this test. The Project library on the otherhand performs roll-backs. This explains its increase in runtime. GHCs library also performs rollbacks and thus its increase in runtime is only natural. We additionally increase the level of concurrency by increasing the TVars that each transaction accesses which leads to more rollbacks in the GHC and Project libraries.

The last three tests use `PerformanceTest`. The number of reads per transaction are equal in all tests, but the `rWRatio` is different. The costs of a transaction are modified slightly, but the level of concurrency is modified formidable. By increasing the number of writes (and preserving the total number of reads), we increase the costs of a transaction slightly, but we increase the level of concurrency noticeable. The more TVars a transaction writes the higher is the chance that another transaction is invalidated. The first test executes ten writes per transaction and ten reads for each write. The results are positive for STMLA since it is the fastest implementation. STMWSL is the slowest implementation. The GHC implementation is slighty slower than STMLA but noticeable faster than the Project implementation. If we compare this to the second test, in which the number of reads remains the same, but

---

[3]We analyze problems and potential performance issues of STMWSL in Appendix B

FIGURE 5.2: Results of the scaling tests with `StmTest`

the number of writes is halved, no considerable changes are observable. All implementations remain on a similar level compared to the first test. If we on the other hand increase the number of writes while retaining the number of reads as in the last test, the runtime of the GHC and Project implementations are increasing. The runtime of STMLA remains similar to the two previous tests. Interesting is that the runtime of STMWSL is lower than before, but still not comparable to the runtime of STMLA. This result is not surprising because the increase in concurrency does not effect STMLA. STMLA does not roll back at all in this test and thus the configuration has no effect on the runtime. The GHC and Project implementation do roll back. Hence a increase in runtime when increasing the level of concurrency is reasonable.

While the results of the first three tests seem not very promising the alternative implementation, the results of the last three tests reveal an application where STMLA outperforms even the GHC implementation which is an omptimized low level `C` implementation. More importantly, these tests met the expectations regarding the costs per transaction. First, the performance does not decrease if we increase the costs per transaction. Since we do not execute any rollbacks in these tests, the costs per transaction do not matter for STMLA (and STMWSL). The GHC and Project libraries on the other hand lose performance in these cases. Second, the performance does not change if we increase the level of concurrency. By increasing the level of concurrency the amount of rollbacks that are executed by the GHC and Project libraries increase. STMLA and STMWSL on the other hand do not suffer from the increased level of concurrency. An explanation and a possible solution for the poor performance of STMWSL is givin in Chapter 6.

### 5.2.2 Second Test Series

We use the second test series to study the scalability of the implementations. We use `StmTest` as well as `PerformanceTest` to compare the four implementations. Similar to the previous test series, the same test is executed multiple times and the average execution time is meassured. The tests are randomized by the test itself and the scheduler. To get a representative result for the runtime, it is inevitable to execute the tests multiple times. To test the scalability we increased the number of

FIGURE 5.3: Results of the scaling tests with `PerformanceTest`

threads over the course of the series. The other configuration parameters remained untouched in the whole test series.

Figure 5.2 shows the result of the scaling test performed with `StmTest`. We used the following configuration: 500 iterations per thread, 100 TVars stored in a list and 20 modifications per transactions with a varying number of threads. The x-axis describes the number of threads and the y-axis contains the total runtime. Note that an increase in threads results in an increased total work load. Thus the increase in runtime in all implementation is mostly due to the increased total workload. All implementations scale equally. The GHC implementation is slightly better and STMWSL is slightly worse than the other implementations. Overall, the runtime increases linear with number of threads. The GHC implementation stands out because of its unstable execution time in this test series. Nevertheless, it still performs better than the other implementations.

The result of the scaling tests performed with `PerformanceTest` is presented in Figure 5.3. The configuration is as follows: 500 transactions per thread, 200 TVar stored in an `IntMap`, five times more reads than write and five writes per transaction (25 read operations per transaction). The results are similar to the previouse results. All implementation scale equally while GHC is a slightly faster and STMWSL slightly slower. The jitter in the execution time of the GHC implementation is less noticeable than in the first test. The execution time of the project implementation on the other hand varies much more in this test than in the first test. Nonetheless, the general tendencies are not different.

This test series show that all implementations are scalable. That GHC performs better than the highlevel implementations is not surprising. In contrast to the tests presented in the previouse section, these test consist of small transactions. The GHC implementation uses a very light weighted locking scheme. This makes the commit phase of the GHC implementation significantly faster than the commit phase of the high level implementations. Thus the fixed (per transaction) overhead of the highlevel implementations is higher than that of the GHC implementation. Since these tests use smaller transactions than the first series, the results are reasonable.

### 5.2.3 Observations

While executing the tests, I noticed that the runtime option `N` that creates multiple OS threads slows down the computation as a whole. The runtime option `N` allows the user to specify the number of OS threads by adding a number. I performed some test with runtime options varying from `N1` to `N4` (since the processor contains four physical cores). The results are clear. The more OS threads the test uses the slower is the test. This is unpleasant because it means our efforts to utilize multi-core processors are futile. To avoid confusion, I am not saying that the computation is less efficient; it is slower. For example, one test takes 1.6 second when executed with `N1` and the **same** test takes 2.3 seconds when executed with `N2` (2.6 with `N3` and 3.6 with `N4`). The test consumed in the case of `N1` a singel core to its full extend (100%). When the test runs with `N2`, it uses two cores, but not to their full extend (around 70% per core). With `N3` it is about 50% per core and with `N4` it is about 45% per core. The system is not utilized by any other process while testing (expect for the OS) . Even though the total amount of CPU consumption increases with the number of OS threads, the execution time as a whole slows down. I did not take any efforts to explore the reason for this. This is most likely a problem in the runtime systems and thus goes beyond the scope of this thesis. The fact remains that in this kind of setup an increase in OS threads results in a decrease in performance. This *kind of setup* is that most of the code is transactional code. Real worlds programs are usually not structured like these tests. The portion of transactional code is far less. Hence this observation does not necessarily mean that STM is useless in practise. Nevertheless, all tests presented in this section are executed with `N4` to allow real parallelism.

In addition to the runtime performance test, I performed tests to check the implementations for memory leaks. The GHC implementation as well as STMLA and STMWSL do not produce memory leaks as far as tested. The Project implementation on the other hand contains a memory leak. When ever a transaction reads a TVar, it enters its `retryMVar` to the queue of the TVar. This allows other transactions to notify it in case they modify the TVar. If the TVar is never written, but continuously read, this queue grows and is never emptied. The transactions themself do not remove the MVar from the queue when they rollback or commit. This is for performance reasons because it requires the transaction to search for its own MVar in all TVar queues it has accessed. In STMWSL and STMLA the only time a transaction adds something to the queue is when it executes `retry`. Additionally, the transaction removes its MVar from the queue when it is no longer needed. In this case the performance is not as important as it is for the Project implementation. The `retry` either suspends the transaction and thus removing the MVars from the queues is a minor performance issue or it does not add the MVar to the queue at all. The other entries of the TVar cannot produce any memory leaks because they only contain one entry at a time. Since the log is discarded each time the transaction rolls back or successfully commits, no unused data is stored within the log.

# Chapter 6

# Conclusion

This chapter contains related work, future work and a summary of the thesis.

## 6.1   Related Work

The most important work for Haskell regarding STM is the paper of (Harris, Marlow, et al., 2005). This work forms the foundation of the current implementation of STM in GHC. Even if the original implementation was reworked over the last years, its core is still present. This core is described in Chapter 2. The modifications of the original library were in most cases bug fixes. The only feature that was added to initial implementation are *invariants* (Harris and Jones, 2006). Besides an implementation sketch and a description of the interface, (Harris, Marlow, et al., 2005) also offer a formal semantics for STM in Haskell. The alternative implementation implements the same interface and aims to fulfill these semantics. There is no formal prove that the alternative implementation (or the GHC implementation) suffice these semantics.

Another approach to optimize STM is presented by (Harris, Plesko, et al., 2006). They propose to use a pessimistic approach comparable to data base transactions. This means each time a shared data structure is accessed it is locked. If the transaction tries to acquire a locked data structure it rolls back. This design avoids *shadow copies* of values such as the logs in the Haskell implementations. This avoids a memory overhead and more important it avoids the necessity to look up values in the log. They made assumptions for the ussage of transactions: Most transactions commit successfully and they use far more reads than writes. Under these assumptions their implementation performs considerable better than optimistic implementations. In the course of this thesis, we did not investigate how a pessimistic implementation performs. Neither made we any assumptions on the usage of STM. The aim was to create a implementation that performs better in general. The timing tests in Chapter 5 show that this was not achieved. There are cases where the alternative implementation performs better than the original implementation and vice versa. In this regard the thesis results are comparable to the results of (Harris, Plesko, et al., 2006). Under the assumption that the transactions are sufficient expensive and the level of concurrency is suffiecent high the new implementation is faster.

In Chapter 2 we defined the current problems of the STM implementation in the GHC. One problem is *when* transactions are rolled back. This problem was solved with this thesis. The other problem is *how* transactions are rolled back. This thesis contains no solution for this problem, but (Agarwal, Gupta, and Kallikote, 2009) present a draft that engages this problem. They present a concept that automatically creates *checkpoints*. If a transaction is rolled back, it is not restarted from the very beginning. The systems validates the checkpoints and restarts the execution from the latest valid checkpoint. The checkpoints basically consist of a log and a list of

remaining operations. In theory, it is useful to create a checkpoint for every read operations on shared data structures. This is not recommended in practice since every creation of a checkpoint consumes time and memory. Thus if the systems creates the maximum amount of checkpoints it never executes operations needlessly multiple times, but its overhead will most likely revokes the performance benefits. The solution is to group action and create checkpoint for these grouped actions. To find an optimal group size remains an open problem.

## 6.2  Future Work

Most of the functions provided by GHC STM are also provided by the alternative implementation presented in this thesis. Nevertheless, there are parts of the external interface and internal semantics that are not covered by the new implementation. This lack of functionality offers the opportunity for further research.

In Chapter 4 the `globalCount` were introduced which is used to create globally unique `ID`s. `ID` is a type synonym for `Int`. In Haskell `Int` is a 32 or 64 bit integer. Thus the number of `ID`s is limited. The only usage of `ID` is to create new TVars. For terminating applications this restriction can be neglected. If the user does not create and delete TVars all the time, he will run out of memory before `globalCount` overflows. If this STM implementation is used within an infinite running systems such as web servers, it may be an issue (assuming that the systems creates and deletes TVars while running). This problem can be avoided by either using an `Integer` which presents an integer of unlimmited size in Haskell or by using a kind of conflict detection to be able to handle the overflow. The first solution adds additional overhead to the implementation, since `Int` is slightly faster than `Integer`. Furthermore, we need to alter the data types used within the implementation. We used `IntMap` for the internal book keeping of the alternative implementation. An `IntMap` cannot handle `Integer`s. The overflow detection can be used in two ways. If the implementation is able to reclaim unused `ID`s, it can reuse these. The conflict detection can also terminate the execution to avoid a miss behaviour of the implementation. Currently the implementation does not detect the overflow and thus may lead to problems in long running systems that use STM.

In Chapter 5 we discovered that STMWSL perfoms poorly in the performance tests. The reasons for this are that the implementation can roll back even if the transaction does not branch. Additionally does the transaction perform two validations to avoid unnecessary locking. The locking is performed by taking the MVars that hold the value for a TVar. By taking this MVar no one is able to read these TVars which is important for the protocol. Unfortunately, the transaction that took these MVars is also not able to read the associated TVars. Thus the transaction must evaluate the delayed reads before it acquires the locks of these TVars. This contains the danger that the TVars are modified after the transaction read them and before it acquires the locks for them. This results in a rollback. If the lock holding transaction were able to read the TVars, it could evaluate the delayed reads after it acquired the locks and validated. On the one hand, this makes the validation less expensive since only the reads that were evaluated in the computation phase are validated. Furthermore, this avoids rollbacks because the values cannot be modified in a manner that it evokes a rollback. After the transaction t1 reads the TVars, it is still possible that another transaction modifies the values that are only read by t1, but not modified. This is not a violation of the ACI properties and thus would not cause a rollback (see (Fraser, 2004)). In conclusion, another locking mechanism that allows the lock holder to read

the locked structure increases the performance of STMWSL significantly. This also adds more complexity to the implementation because reading in a locked context contains the risk of deadlocks. Avoiding these deadlocks adds additional complexity to the implementation.

The alternative implementation lacks a sufficient support for exception handling. If an exception is raised, this exception is thrown to the caller. This behaviour is not incorrect per se. If the exception is raised because the transaction has seen an inconsistent view of the memory, it is incorrect. When the transaction raises an exception during the execution, the implementation must check whether the transaction is valid or invalid. If it is invalid the transaction have to be rolled back, otherwise it is a violation of the consistency property. If the transaction is valid, the behaviour is a design decision. Either the transaction is raised to the caller or the transaction handles the exception like a retry, meaning it waits for a change of the read TVars and restarts. Both execution paths preserve the ACI properties and thus are legit. There are additional technical problems with exceptions in the implementation. If an exception is raised during the commit phase, no measures are performed to make sure the system is still usable. If the transaction already locked several TVars and then raises an exception, this exception is not catched and the locks for the TVars are not released. Thus it is possible that the TVars remain locked. I am not sure if it is possible that an exception is raised at this point because only read and write operations on MVars and safe operations on the `IntMaps` are performed. Nevertheless, it remains an unsolved problem.

(Harris and Jones, 2006) introduce a new feature to STM called *invariants*. This allows the user to define invariants on the TVars that must be fulfilled before a transaction can commit. If a transaction tries to commit a state that violates at least one invariant the transaction is rolled back (or suspended in case none of the read TVars has changed). This allows the user to compose STM functions even better for the costs of performance. An integration of invariants in the alternative implementation inceases the usability of the alternative implementation to a level that is similar to GHC STM.

STMLA has proven to be competitor to the GHC implementation of STM with regards to performance, even though the implementation makes no use of low level C primitives to boost its performance. There are cases where the GHC implementation performs better. To implement STMLA with C primitives and integrating it in the compiler provides multiple benefits. The performance increases could be enough to substitue the current implementation in GHC. This need to be investigated with additional performance test. Like stated in Chapter 5, this is a difficult topic, since STM is a universal tool and cannot be tested easily. Even if the implementation is not fast enough to replace the current implementation, it can provide an alternative to the current implementation which is used for specific cases. Then the user (or compiler) can decide which implementation he wants to use. More importantly the integration in the compiler allows the implementation to detect conflict before the commit phase. This is important if the transaction executes an infinite loop because it saw an inconsistent view of the memory. Without an early conflict detection the transaction does not break the loop because it never validates before the commit phase. The GHC implementation uses the runtime system to evoke validations in the computation phase to avoid these kinds of loops. A compiler integration of the alternative implementation allows a similar mechanism.

We introduced in Section 2.3.2 another problem of the GHC implementation [1]:

---

[1]This Problem is also present in the Project implementation and the alternative implementations.

the way transactions are rolled back. Regardless the reason for a rollback, the transaction starts all over. This includes that the transaction executes actions on TVars that have not changed. Unfortunately, I was not able to implement a suitable solution for this problem in the time of this thesis. The previously presented paper (Agarwal, Gupta, and Kallikote, 2009) had some success engaging this problem. To provide a similar solution for Haskell may be possible by redefining the monadic operations » and »=. The first step is to divide the transaction in the parts that are seperated by » or »= (the single *statements* of the do-notation). We call these parts chunks for now. The second step is to deteremine which TVars are used in which chunks. `unsafePerformIO` can be used to do this. » and »= both execute the first action. During this execution `unsafePerformIO` can be used to detect the values (and the associated TVars) that are needed to evaluate this chunk. In the end the transaction contains an ordered collection of pairs. Each pair consist of a chunk and a set of TVars that this chunk depends on. In the case the transaction is invalid it can be rolled back to the newest chunk that is still valid. This includes that every time the transaction logs a chunk and its dependencies, the transaction also needs to backup the read/writeLog. Thus this implementation contains a lot of book keeping. It needs further research to decide whether this idea works correctly and if so whether it increases the performance despite the book keeping overhead.

## 6.3   Summary

In Chapter 1 we gave a motivation for synchronization in general. To utilize multi-core processors and guarantee reactivity is mandatory for software nowadays. Multithreading is essential to achieve these objectives. Multithreading also includes dangers such as lost updates or race conditions. These problems are engaged with synchronization tools, starting from semaphores to abstract concepts such as STM.

Chapter 2 first motivated STM in Haskell by comparing it to MVars. The usage of TVars and MVar is similar, but TVars or more specifically STM guarantee the ACI properties which makes deadlock avoidance much easier to handle. Furthermore, we explored the current implementation of STM in GHC. This library makes use of low level C primitives and the runtime systems to ensure its correct behaviour and performance. At last, we examined the performance problems in this implementation and introduced the idea of delaying IO-reads to avoid one problem. The problems are *how* and *when* a transaction is rolled back. This also definied the aims of the thesis, i.e. to provide an implementation that avoids these problems.

The concept to provide an implementation that avoids one of these problems is presented Chapter 3. First we identified the technical reason for unnecessary rollbacks. We introduced the term *critical TVars* to denote TVars whos modification results in a rollback. Furthermore, we specified when it is mandatory to evaluate an IO-read; if the value is used in a branch condition. The key idea is to minimize the time the TVars are critical. This is achieved by delaying the evaluation of IO-reads as far as possible. The result is that IO-reads, that are not needed for branch conditions, are delayed to the commit phase and others are evaluated just before they are used.

The main part of this thesis is presented in Chapter 4. The implementation of the previously presented idea is described in detail in this Chapter. The STM monad is a state monad and processing a transaction is devided in two phases, the computation phase and the commit phase. In the computation phase the state is enriched by the `writeTVar` and `readTVar` operations. In the commit phase this state is used to validate the transaction an possibly publish its results. To delay the execution of the

IO-read and to find out when a value is needed for a branch condition, the values are wrapped by an `unsafePerformIO` action. If this action is not processed in the computation phase, its execution is forced in the commit phase to ensure the reads are evaluated before the transaction finishes.

The new implementation is compared to the original implementation and a reference implementation in Chapter 5. The results of these performance tests look promising. The new implementation is in all tests faster than the reference implementation and in some test even faster than the GHC implementation. Nevertheless, more tests are needed to verify if the implementation is faster in general. The expectations regarding the execution time, when increasing the level of concurrency, are also met. If we increase the level of concurrency and no TVars are critical, the alternative implementation does not lose performance.

**Conclusion**

The aim of this thesis was to provide an alternative implementation of STM in Haskell to avoid the problems of unnecessary rollbacks and unnecessary recomputations. The presented implementation is a pure Haskell implementation in contrast to the original implementation that uses C primitives. The alternative implementation uses `unsafePerformIO` to avoids the unnecessary rollbacks. The previous section discussed an idea to solve the problem of unnecessary recomputations. However, this idea is not implemented and tested in the alternative implementation. The test results have shown that the new implementation performs better than a reference implementation and equally to the original implementation. Besides the missing support for exception handling and invariants, the alternaive implementation provides the same functions as the original implementation.

# Appendix A

# Applicative

The thesis started with the observation that the rollback explained in section 2.3.1 is unnecessary. It is a modify operation that does not depend on the actual value of the TVar. Thus read the value, modify it and write it back is a imperative way of solving this task. In other words, it is unnecessary to read the TVar during the computation phase, when the transaction does not branch on that value. At this point we did not see that it is not needed when its evaluation is not demanded. In order to avoid the rollback, we wanted a mechanism that prohibits the user from branching on the value. The user should still be able to use that value to do pure computations and write it to TVars. This allows the library to decide when the IO-read is evaluated. We decided to use Applicative[1]. Applicative is comparable to Monad. Additionally did we change the type of `writeTVar :: TVar a → STM a → STM ()`. This allows the following example:

```
inc tv = writeTVar tv $
         pure (+1) <*>
         readTVar tv
```

We are able to modify the value of a TVar without »=. Furthermore did we changed the interal semantics, that an IO-read is performed in either »= or in the commit phase. It is not performed in `readTVar`. Thus if we read a TVar and use applicative operations instead of monadic operations, we do not risk a rollback. The value is evaluated in the commit phase, hence it is not critical at all. This approach avoids the intially mentioned rollback, but has its drawbacks.

The use of applicative operations is not as comfortable as monadic operations. One problem is the use of multiple operators instead of just one. There is no suitable language feature such as the do-notation like in the monadic version. We cannot use `ApplicativeDo` because the target of `ApplicativeDo` is something else[2]. Consequently, we need to use brackets or `$` to group the actions. Another problem is the reversed order of operations. In the monadic version, the operations are ordered like in an imperative language, for example: read → (+ 1) → write. This is easily comprehensible. The applicative version uses a functional style, for example: write (1 + (read)). For this example it is comprehensible, but you can imagine what happens if the actions increase in size. For every operation we introduce in the applicative version another nesting. This leads very fast to incomprehensible code, espeacially when using functions with more than one argument. The monadic version is just a chain of operation, which can easily be extended.

---

[1]https://hackage.haskell.org/package/base-4.9.0.0/docs/Control-Applicative.html

[2]`ApplicativeDo`s aim is it to use applicative operations to calculate an result and not to process a chain of operations. To understand the problem, we would need to understand the translation scheme of `ApplicativeDo` which goes beyond the scope of this thesis.

In conclusion we gain performance for the costs of usability of the library; One of the biggest strength of STM. We solved one problem by introducing another problem. Nevertheless, I implemented a version of STM that uses this scheme and avoided the targeted rollbacks. `IO ()` actions were used to delay the evaluation[3]. When »= is executed these `IO` actions are evaluated. The `IO` actions are created by `readTVar`. Remember that `readTVar` creates an STM action. This action can either be used to bind it with »= or it can be used to write another TVar (the value may be modified after read and before written). If it is wrtten to a TVar the `IO` action was logged and executed in the commit phase. This solution contained some problems.

The question arised how we are able to execute and action and use it at the same time. We need to execute it to make sure the correct value is the result of the action. This can either be the actual value of the TVar in a form of an `IO` action or it is an `IO` action that is written in the log. The following example would not work:

```
swap t1 t2 = do
    let a = readTVar t1
    writeTVar t1 <*> readTVar t2
    writeTVar t2 a
```

This swap function does not work, because the `readTVar` operation in the let expression is evaluated, when `writeTVar t2 a` is evaluated. At this point the value of `t1` is no longer ins initial value, but the value of `t2`. Hence in the end of the transaction both TVars contain the same value. To solve this problem we introduced the function `eval :: STM a -> STM (STM a)`. This funcion allows us to evaluate the an `STM` action an us it at the same time. It first evaluates the `STM` action and then returns an dummy action that has no side effects and returns the value the intial action would return at the point `eval` was called. For example:

```
swap t1 t2 = do
    act <- eval $ readTVar t1
    writeTVar t1 <*> readTVar t2
    writeTVar t2 act
```

This swap function is correct. The first action evaluates `readTVar t1` and creates an action that return the same as `readTVar t1` at this point. This works for every `STM` action. Side effect evoked by the action on the other hand are not evoked by the returned action. Side effect in the context `STM` are always modifications of the `StmState`. However, the `eval` function also solved another problem. The following example is not possible without eval:

```
double t1 t2 t3 = do
    act <- eval $ readTVar t1
    writeTVar t2 act
    writeTVar t3 act
```

It is not possible to write this example without evaluating the `readTVar t1` twice or using `eval`. `eval` can be compared to `share`[4] whichs allows explicit sharing for action results.

Even though `eval` allows us to use a readTVar action twice, there are still problems. The previous example would creates log entries for `t2` and `t3` that both yield the same `IO` action. The `IO` action that IO-reads `t1`. In the commit phase these action are evaluated. This includes that the actual TVar is read twice because the action is

---

[3]This implementation did not use `unsafePerformIO`. It used monadic `IO` actions.
[4]https://hackage.haskell.org/package/explicit-sharing-0.9/docs/Control-Monad-Sharing.html

present in two log entries. At the first glance it does not seem to be an issue, but also pure computation are duplicated as the following example emphasises:

```
shareless t1 t2 t3 = do
  act <- eval $ pure sqrt <*> readTVar t1
  writeTVar t2 act
  writeTVar t3 act
```

In this example the pure function `sqrt` is also evaluated twice. This is necessary in general because the IO-read (created by readTVar) could return different values depending on the time it is executed. For our case we know that this is not the case. Both IO-reads are performed in the commit phase. At this time the TVar `t1` is locked. Thus the actual value of `t1` cannot change (as long as we do not change it ourself).

The core problem is that Haskell does not allow sharing on `IO` actions dirrectly. There is a way to enable sharing on `IO` actions: `unsafePerformIO`. These actions are considered as pure values, hence their results are shared as any other pure value in Haskell. While using using `unsafePerformIO` to implement the sharing in the applicative solution, we discovered that it can replace the whole need for Applicative. Since the usability of the applicative solution was still worse than the usability of the original implementation, we discarded the applicative implementation and created the implementation that is presented in the main part of this thesis.

# Appendix B

# STMWSL

In the last days of writing this thesis I reviewed the code of STMWSL and found multiple performance issues and a bug that involves the potential for deadlocks. Due to the limited time, I was not able to solve these problems and fix the bugs. In this appendix we will investigate these problems and the bug as possible reference for future work.

The general scheme of STMWSL in the commit phase is the following:

1. validate

2. evaluate reads

3. lock written TVars

4. validate

5. commit

STMWSL processes the writeSet multiple times to achieve this scheme. When locking the written TVars the implementation needs to find out which TVars were written and thus processes the whole writeSet. After that the implementation locks the TVars that are written. Additionally are these TVars validated in case they were read. This includes that the readSet (see Section 4.1) is processed. The validation of these TVars is needed at this point because after the TVars are locked no furhter reading (and hence validation) is possible. This is the first performance issue. It is not needed to devide this into two phases. The locks for the TVars could be acquired at the same time the writeSet is processed to find out which TVars need to be locked. Another solution to this problem is to refrain from explicitly evaluating the unevaluated reads. By processing the writeSet all reads that are needed are evaluated (all but the ones that are needed for the result of the transaction). In other words, the current implementation needs for the second and third part about (`2*read + accessed + written`) operations. `read` denotes the number of different TVars that were read during the transaction, `accessed` is the size of the writeSet and `written` is the number of different TVars that were written. With the proposed solutions the operations that are needed are either (`2*reads + accessed`) or (`reads + accessed + written`).

To be able to validate, the implementation uses IORefs. A direct value comparison is not suitable in Haskell. The problem is Haskells non-strict semantics. In general values are not evaluated, at least not completely evaluated. Thus by comparing these values we need to evaluate them. This on the other hand changes the semantics of our program. The solution to this problem is to wrap every value in an IORef. These IORefs are used as immutable variables, meaning whenever we change the value, we also create a new IORef. This allows us to compare the IORefs

to validate and the value itself remains unevaluated. The drawback is that we need to create new IORefs every time we write a TVar. This is a significant overhead that the original implementation does not has. The original implementation uses its compiler integration and compares the values directly without evaluating them. Haskell values are represented as pointers. These pointers can be accessed and compared by the GHC implementation thanks to their use of C primitives and compiler integration. In a high level Haskell library the only way to compare the pointer of value is by using `reallyUnsafePtrEquality#`. As the name suggests the results of this function may vary, because pointer in Haskell have a very limited validity. The garbage collector sometimes rearrange values in the head. This also changes the pointers of the values.

Besides these performance issues, there is a bug is the implementation. The problem is the second validation. This implementation uses a locking mechanism that prevents all threads from reading the TVars once they are locked. On the other hand it allows us to only lock the TVars that are written instead of all TVars that were accessed. To understand the problem we need to take lock at the following example. Thread 1:                                         Thread 2:

```
transaction1 = do
  a <- readTVar t1
  writeTVar t2 a
```

```
transaction2 = do
  a <- readTVar t2
  writeTVar t1 a
```

In STMWSL the following execution is possible. `Thread1` executes until it has acquired the lock for `t2`. Then `Thread 2` executes until it has acquired the lock for `t1`. Thus both threads finished the third phase of the commit scheme. If now either of the Threads try to validate, it needs to read a locked TVar (the TVar locked by the other thread). This results in the threads suspension. This shows the potential for deadlocks in this implementation. Fortunately, in the current implementation is another bug that prevents the second validation to work and thus this deadlocks did never occurred. Unfortunately this also violates the ACI properties. Small tests have schown that it is possible to cause a lost update. Before validating the second time, we need to remove the locked TVars from the readSet to avoid that a transaction deadlocks itself. For example if the transaction locked TVar `t1` it does not need to validate it because locking always includes validating. On the other hand if we try to validate the TVar after it is locked, the transaction deadlocks. The before validating the second time we used `IntMap` operations to remove all entries from the readSet that are also in the writeSet. The readSet is a subset of the writeSet and thus the validation was performed on the empty readSet. To fix this bug requires to identify the written TVars and remove them from the readSet. This on the other hand require the transaction to process the writeSet once more and thus is a significant overhead.

These obeservations suggest that a change in the `StmState`. We need an efficient way to access the written TVars, which does not require to process the whole writeSet. To avoid the deadlock on the other hand, requires either recursive helping as suggested in (Fraser, 2004) or the transaction needs to rollback if it tries to access a locked TVar are in the original implementation.

# Appendix C

# Source Code

All outputs of this thesis are stored in the GitHub `https://github.com/LasFo/ma-the`.

The structure is as follows:

- `tests` contains the tests that were developed during the thesis. Most of thes test were used to test for specific bugs during the development process. Thus some of these test seem unnecessary. The subdirectory contains tests that run indefinitely.

- `implementation` contains the implementations STMWSL, STMLA and the Project implementation. It does not contain the GHC implementation. Since it is no longer part of the standard libraries of GHC 8.0.1, it can be downloaded with `cabal`[1]

- `recavoid` contains ideas and challenges for the problem of unnecessary recomputations.

- `stm` contains all code that were developed during the the first part of this thesis that is described in Appendix A.

- `timing` contains the ressources that were used for the performance tests.

- `others` contains a documentation, this written part and slides for a talk.

The implementation was developed and tested with GHC 8.0.1. Other version were not tested. There is no guarantee that the implementation works for other versions. Serveral libraries are used within the implementation. All used libraries are part of the standard Distribution of GHC. The only library that is used but not in the standard Distribution is `Data.Tuple.Select`. There cabal package `tuple` includes this library.

---

[1] `https://www.haskell.org/cabal/`

# Bibliography

Agarwal, Shivali, Monika Gupta, and Shyamasundar Rudrapatna Kallikote (2009). *Automatic checkpointing and partial rollback in software transaction memory*.

Dijkstra, Edsger Wybe (1965). "Cooperating Sequential Processes, Technical Report EWD-123". In:

Fraser, Keir (2004). *Practical lock-freedom*. Tech. rep. University of Cambridge, Computer Laboratory.

Gray, Jim and Andreas Reuter (1992). *Transaction processing: concepts and techniques*. Elsevier.

Harris, Tim and Simon Peyton Jones (2006). "Transactional memory with data invariants". In: *First ACM SIGPLAN Workshop on Languages, Compilers, and Hardware Support for Transactional Computing (TRANSACT'06), Ottowa*. Vol. 92.

Harris, Tim, Simon Marlow, et al. (2005). "Composable Memory Transactions". In: *Proceedings of the Tenth ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*. PPoPP '05. Chicago, IL, USA: ACM, pp. 48–60. ISBN: 1-59593-080-9. DOI: 10.1145/1065944.1065952. URL: http://doi.acm.org/10.1145/1065944.1065952.

Harris, Tim, Mark Plesko, et al. (2006). "Optimizing memory transactions". In: *ACM SIGPLAN Notices* 41.6, pp. 14–25.

Huch, Frank and Frank Kupke (2005). "A high-level implementation of composable memory transactions in concurrent haskell". In: *Proceedings of the 17th international conference on Implementation and Application of Functional Languages*, pp. 124–141.

Larus, James R and Ravi Rajwar (2007). *Transactional memory*. Vol. 1. 1. Morgan & Claypool Publishers.