

Software Transaction Roll Back Avoidance

Master Proposal Talk

Three horizontal bars of different shades of green and blue, stacked on top of each other, spanning the width of the slide.

Lasse Folger

November 25, 2016

Motivation

```
type Account = MVar Int
```

```
transfer :: Account -> Account -> Int -> IO ()
```

```
transfer src dst am = do
```

```
  a1 <- takeMVar src
```

```
  a2 <- takeMVar dst
```

```
  putMVar src (a1 - am)
```

```
  putMVar dst (a2 + am)
```

MVar

Thread 1:

```
transfer acc1 acc2 50
```

Thread 2:

```
transfer acc2 acc1 50
```

MVar

Thread 1:

```
a1 <- takeMVar acc1
a2 <- takeMVar acc2
writeMVar acc1 (a1 - 50)
writeMVar acc2 (a2 + 50)
```

Thread 2:

```
a1 <- takeMVar acc2
a2 <- takeMVar acc1
writeMVar acc2 (a1 - 50)
writeMVar acc1 (a2 + 50)
```

MVar

Thread 1:

```
a1 <- takeMVar acc1
a2 <- takeMVar acc2
writeMVar acc1 (a1 - 50)
writeMVar acc2 (a2 + 50)
```

Thread 2:

```
a1 <- takeMVar acc2
a2 <- takeMVar acc1
writeMVar acc2 (a1 - 50)
writeMVar acc1 (a2 + 50)
```

⇒ Deadlock

Use Transactions

```
type Account = TVar Int

transfer :: Account -> Account -> Int -> STM ()
transfer src dst am = do
  a1 <- readTVar src
  a2 <- readTVar dst
  writeTVar src (a1 - am)
  writeTVar dst (a2 + am)
```

TVar

Thread 1:

```
atomically $  
  transfer acc1 acc2 50
```

Thread 2:

```
atomically $  
  transfer acc2 acc1 50
```

TVar

Thread 1:

```
atomically $  
  transfer acc1 acc2 50
```

Thread 2:

```
atomically $  
  transfer acc2 acc1 50
```

⇒ works fine, because transactions provide ACID properties

Current Implementation (Control.Concurrent.STM)

- *writeTVar*, *readTVar* and *newTVar* modify TVars
- *retry* and *orElse* alter the control flow
- *atomically* executes a transaction
- composition via bind operator (or do)

Modify Operations

- **newTVar**: creates a new, initialized TVar
- **readTVar**: add the TVar and its version number to the read set
- **writeTVar**: add the TVar and the value to the write set
- read set and write set serve as local cache

atomically $:: STM\ a \rightarrow IO\ a$

1. compute the read set and write set
2. validate the read set
3. if valid commit the write set
4. else restart

Validation

1. compare version number in read set with actual version number
2. if there is a difference return invalid
3. else return valid

Problem

Thread 1:

```
a1 <- readTVar acc2  
a2 <- readTVar acc1  
writeTVar acc2 (a1 - 50)  
writeTVar acc1 (a2 + 50)
```

Thread 2:

```
a1 <- readTVar acc2  
a2 <- readTVar acc1  
writeTVar acc2 (a1 - 50)  
writeTVar acc1 (a2 + 50)
```

Problem

Thread 1:

```
a1 <- readTVar acc2
a2 <- readTVar acc1
writeTVar acc2 (a1 - 50)
writeTVar acc1 (a2 + 50)
```

Thread 2:

```
a1 <- readTVar acc2
a2 <- readTVar acc1
writeTVar acc2 (a1 - 50)
writeTVar acc1 (a2 + 50)
```

⇒ either sequential or one transaction is rolled back

Idea

Thread 1:

```
a1 <- readTVar acc1  
writeTVar acc1 (a1 - 50)  
a2 <- readTVar acc2  
writeTVar acc2 (a2 + 50)
```

Thread 2:

```
a1 <- readTVar acc2  
writeTVar acc2 (a1 - 50)  
a2 <- readTVar acc1  
writeTVar acc1 (a2 + 50)
```

Idea

Thread 1:

```
a1 <- readTVar acc1
writeTVar acc1 (a1 - 50)
a2 <- readTVar acc2
writeTVar acc2 (a2 + 50)
```

Thread 2:

```
a1 <- readTVar acc2
writeTVar acc2 (a1 - 50)
a2 <- readTVar acc1
writeTVar acc1 (a2 + 50)
```

⇒ delay the evaluation of readTVar to avoid rollback, but...

Idea does not work

```
limitedTransfer src dst am = do
  a1 <- readTVar src
  if a1 < am
    then return ()
    else do a2 <- readTVar dst
             writeTVar src (a1 - am)
             writeTVar dst (a2 + am)
```

... the transaction may need the value

Problem

- $(>>=) :: \text{STM } a \rightarrow (a \rightarrow \text{STM } b) \rightarrow \text{STM } b$
- Bind extracts the value from the STM context
- STM no longer controls this value.
- Need another Typeclass than Monad

Applicative

- Applicative is less powerfull than Monad
- $(\langle * \rangle) :: \text{STM } (a \rightarrow b) \rightarrow \text{STM } a \rightarrow \text{STM } b$
- The value can be modified without leaving the STM context

Project

- Improved a pure Haskell implementation
- Direct notification
- Explicit, ordered locking
- Optimisations

Master Thesis until now

- Composition by combination of Monad and Applicative
- $\gg=$ evaluates and enables rollbacks
- $\langle * \rangle$ enables to modify values without rollbacks
- `writeTVar :: TVar a -> STM a -> STM ()`

New Combinators

- $(\langle * \rangle) :: \text{STM } (a \rightarrow b) \rightarrow \text{STM } a \rightarrow \text{STM } b$
- $(\langle ** \rangle) :: \text{STM } a \rightarrow \text{STM } (a \rightarrow b) \rightarrow \text{STM } b$
- $(\langle * \rangle) :: \text{STM } a \rightarrow \text{STM } b \rightarrow \text{STM } b$
- $(\langle ** \rangle) :: \text{STM } a \rightarrow (\text{STM } a \rightarrow \text{STM } b) \rightarrow \text{STM } b$
- $(\langle > > = \rangle) :: \text{STM } a \rightarrow (a \rightarrow \text{STM } b) \rightarrow \text{STM } b$
- $(\langle > > \rangle) :: \text{STM } a \rightarrow \text{STM } b \rightarrow \text{STM } b$

New Transfer

```
transfer :: Account -> Account -> Int -> STM ()
transfer src dst am = do
    readTVar src <*> pure (- am) *> writeTVar src
    readTVar dst <*> pure (+ am) *> writeTVar dst
```

Problem solved

Thread 1:

```
atomically $  
  transfer acc1 acc2 50
```

Thread 2:

```
atomically $  
  transfer acc2 acc1 50
```

⇒ no more rollback

Todo

- Reduce the number of combinators
- ApplicativeDo might do that
- investigate other problems:
 - Branch condition is not changed by TVar modification
 - Recomputation of values which did not change

Unnecessary Recomputation

```
transaction = do  
  limitedTransfer acc1 acc2 50  
  limitedTransfer acc3 acc4 100
```

If one transfer is invalidated, both are recomputed

Questions about...

- ...Control.Concurrent.STM?
- ...rollback avoidance?
- ...unnecessary recomputation?
- ...something else?